

Watermarking (Defense and Hazards)

Luca Caviglione

Institute for Applied Mathematics and Information Technologies
National Research Council of Italy
luca.caviglione@ge.imati.cnr.it

Summer School on Artificial Intelligence for a Secure Society

5 - 10 September 2024

Capo Vaticano Resort Thalasso SPA
Capo Vaticano, Calabria, Italy



Consiglio Nazionale
delle Ricerche

Course Goal and Organization

- This is a **crash course** on **watermarking**
- Spoilers:
 - techniques are tightly-coupled with the application
 - emphasis on the AI world
 - watermarking is a double-edged sword!
- Organization
 - Day 1 (9 September – 1h): basic information, the big picture, and examples
 - Day 2 (10 September – 2h): more examples, security, and challenges
- All the needed material is available on the course GitHub:
 - <https://github.com/lucacav/AI4SS>



Motivations (1)

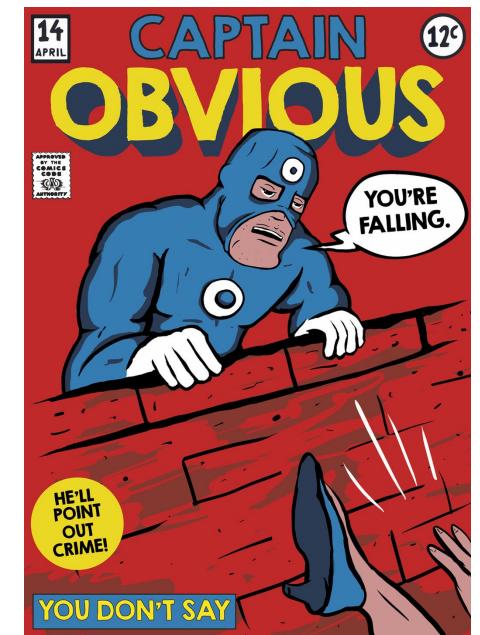
- **Explosion** in the diffusion of **digital information**
- Some examples:
 - **digital media**: audio, video, pictures, and text
 - **software**: source code, binaries, and libraries
 - **3D mesh**: entertainment, lidar scans, 3D-printed artifacts
 - **urban intelligence applications**: IoT data, maps, and measurements
 - **medical data**: imaging and ad-hoc prostheses
 - **network traffic**: routing, QoS, and traffic engineering
 - **hardware**: FPGA and custom silicon
 - **artificial intelligence**: datasets, models, and outputs

Motivations (2)

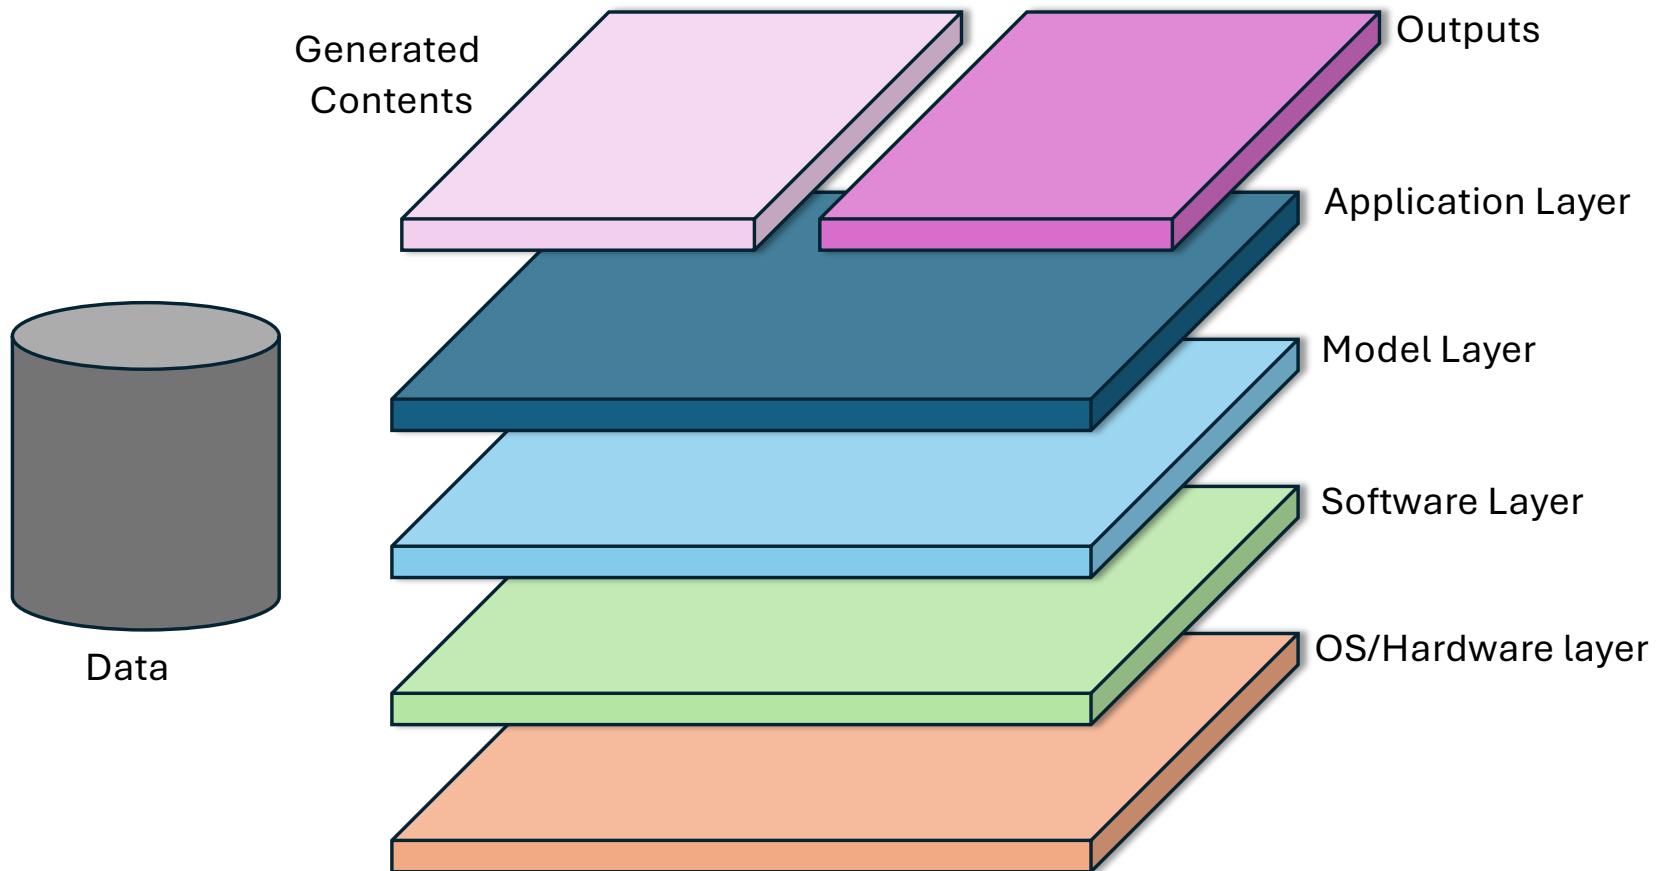
- Technological, ethical, and economic **drivers**:
 - diffusion of the “as-a-Service” paradigm
 - industrial competition
 - digital sovereignty
 - protection of the value
 - misinformation and manipulation
- **Digital information** requires (efficient) **mechanisms** for:
 - copyright protection
 - fingerprinting and tracking
 - integrity
 - authentication
 - annotation/metadata

Watermarking and AI (1)

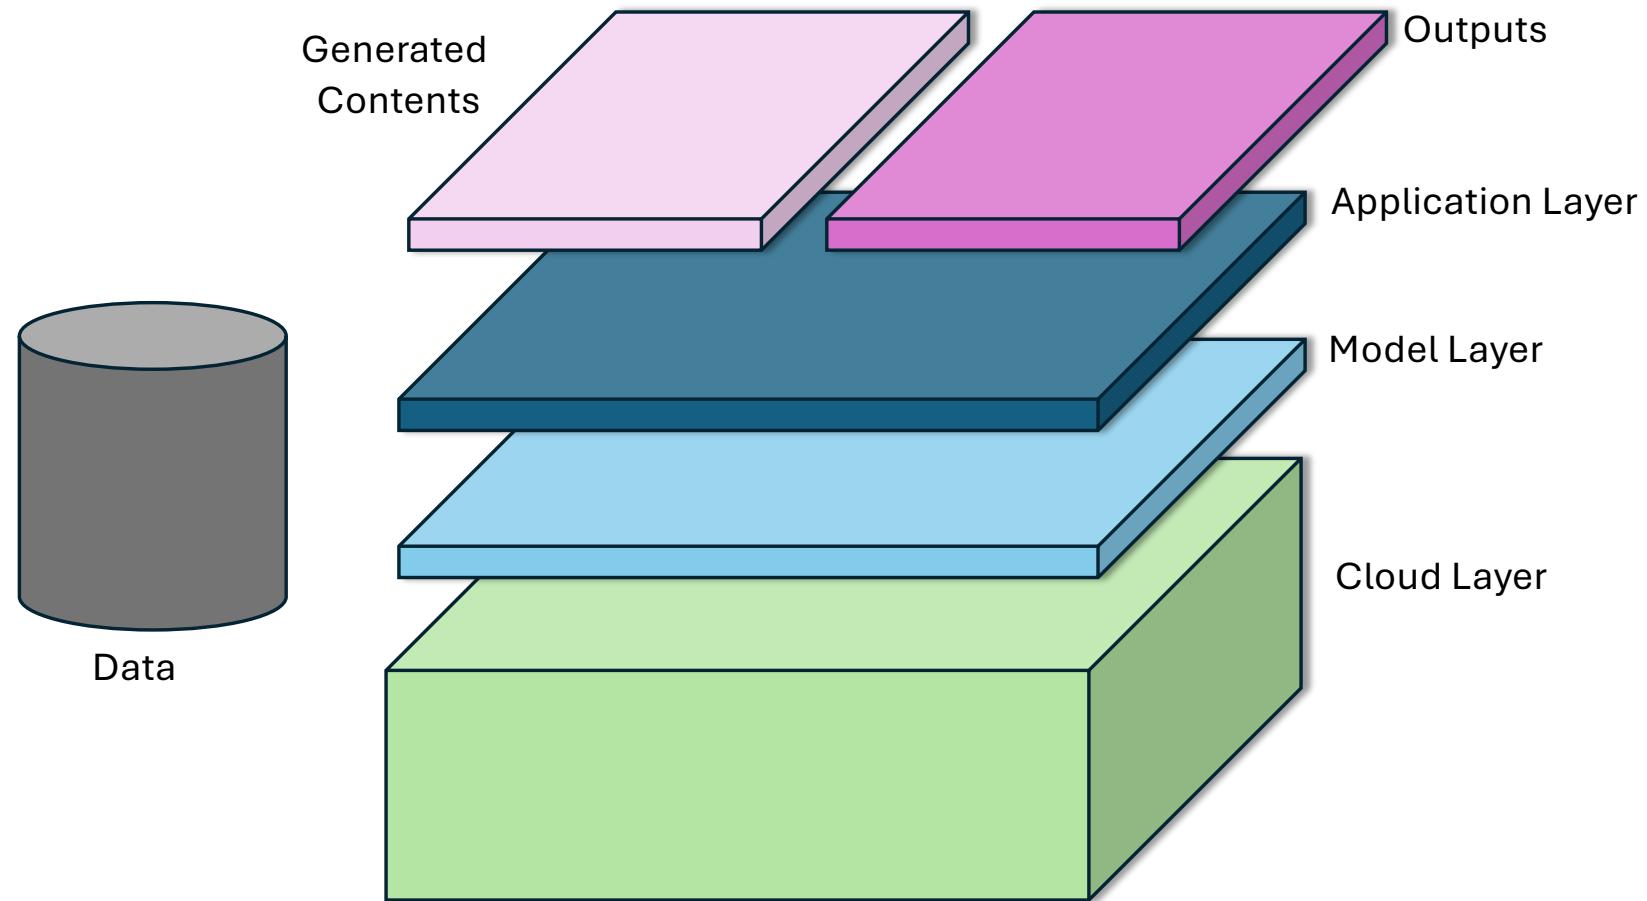
- **Watermarking** is about **embedding** information into original data to:
 - claim something, e.g., copyright/ownership
 - not affect the usability
- AI is **not** only about AI:
 - datasets and (pre-trained) models
 - software libraries and running code
 - dedicated hardware
 - generated contents and as-a-Service frameworks
- There are many **issues** that (still) require a **solution**:
 - detect poisoning
 - enforce guarantees in ML-as-a-Service
 - copyright
 - ...



Watermarking and AI (2)



Watermarking and AI (2)

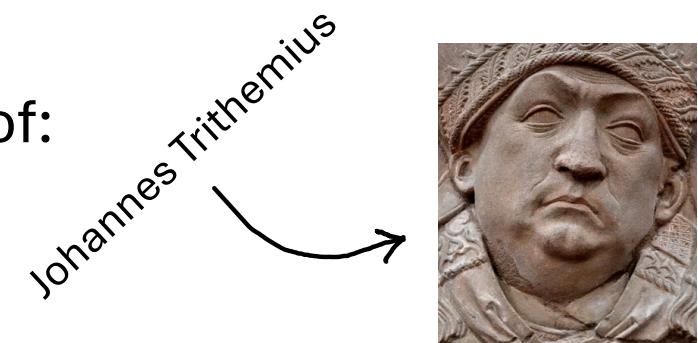


Watermarking and AI (2)

There is no a one-size-fits-all mechanism for watermarking all the entities composing an AI stack!

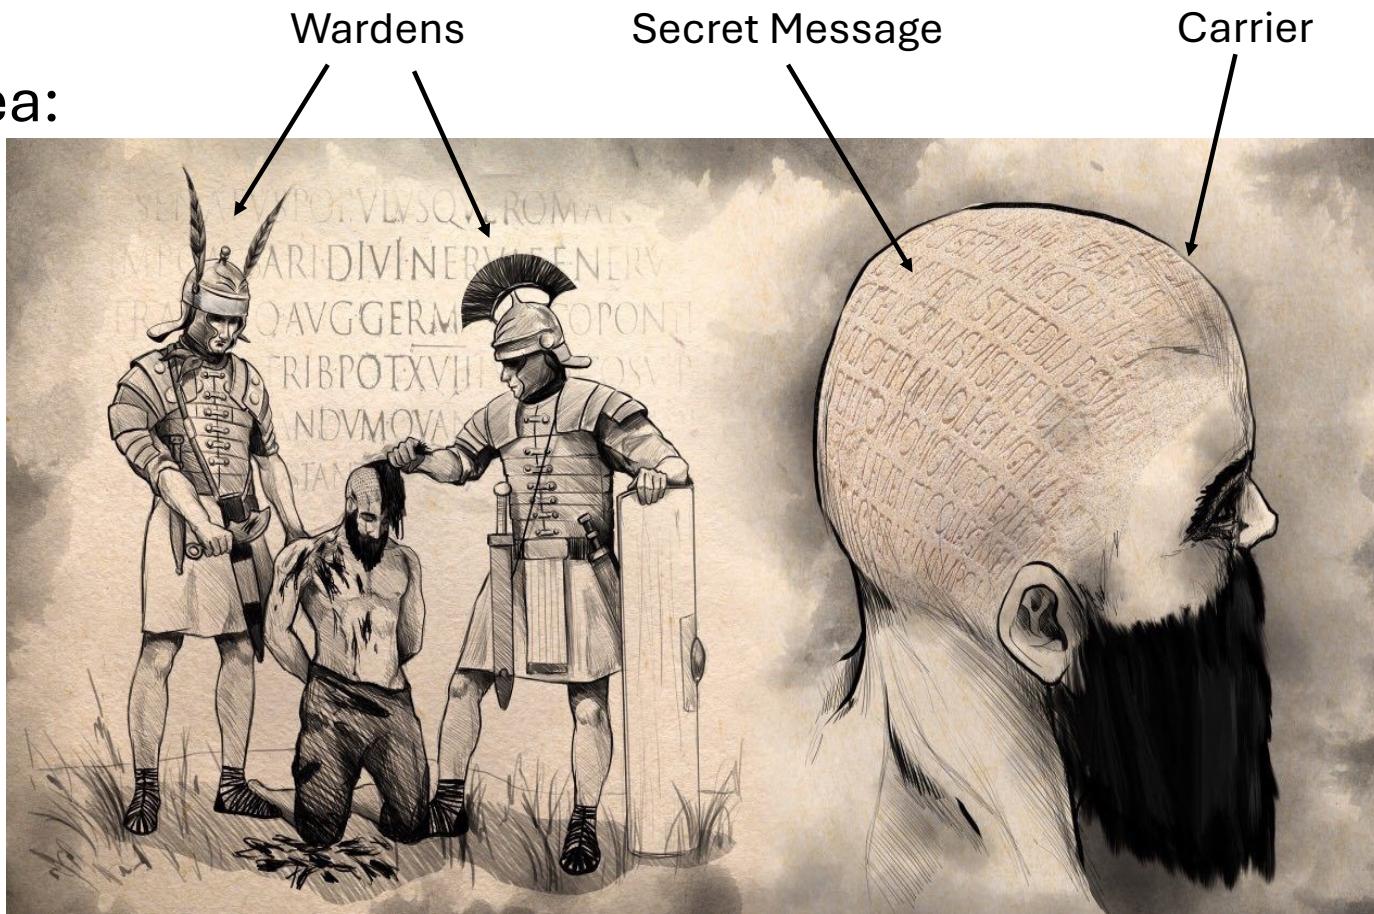
Basic (1)

- The word **steganography** is the combination of:
 - *steganos* (στεγανός) = covered or concealed
 - *graphe* (γραφή) = writing
- The **first** recorded use of the term:
 - in 1499 by Johannes Trithemius
 - book “Steganographia”, i.e., an essay on cryptography and steganography
- Mentioned in 440 BC by Herodotus in his Histories
- **Cloak secret data into a suitable carrier**



Basic (2)

- Key idea:



Source: <https://medium.com/@z3roTrust/using-digital-steganography-to-protect-national-security-information-463bba664830>

Basic (3)

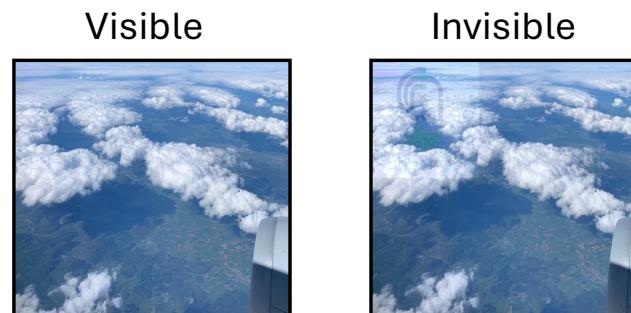
- Watermarking mechanisms have many **properties/requirements**
- Basic:
 - **capacity**: how much information can be embedded
 - **robustness**: how the data can withstand (un)wanted removal attempts
 - **secrecy**: how difficult is to spot the secret information
- Other(*):
 - **fidelity**: how much degradation is caused to the carrier
 - **reliability**: owner can identify secrets with a high probability, i.e., low F_N
 - **integrity**: owner rarely accuses another owner, i.e., low false alarm
 - **efficiency**: the price paid for embedding/verifying the watermark
 - **generality**: is the mechanism dependent on a specific carrier?

(*) F. Boenisch, "A Systematic Review on Model Watermarking for Neural Networks", Frontiers in Big Data, Vol. 4, pp. 4:1 – 4:16, Nov. 2021

Basic (4)

- Another possible characterization:

- **visible**
- **invisible**

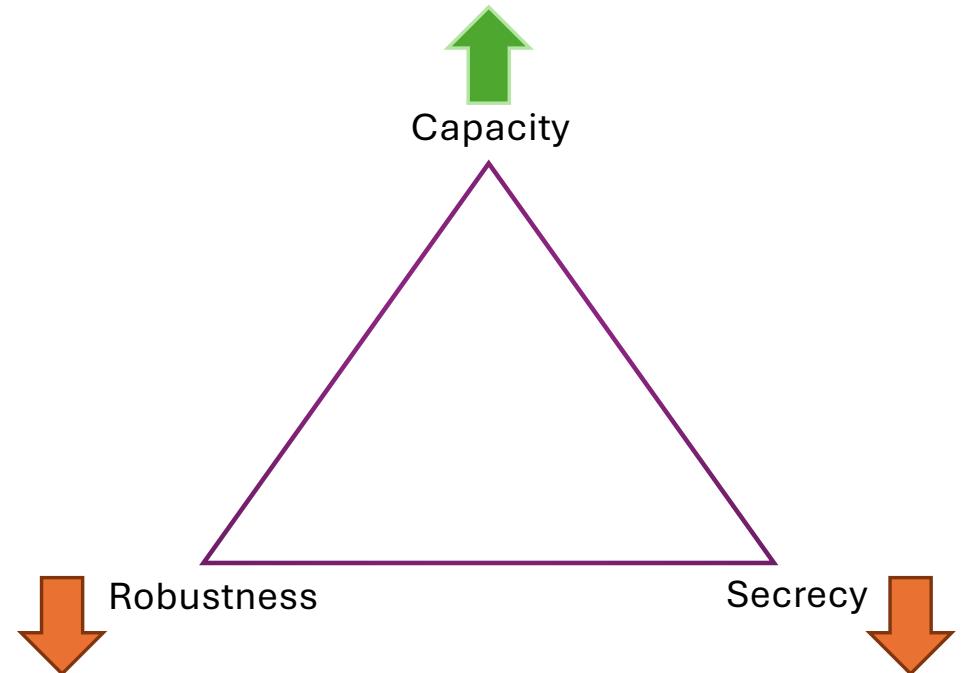


- The ultimate **goal** of the watermark also plays a role:

- **robust**: suitable if the carrier is expected to be processed multiple times (e.g., transcoding) or attacked by a malicious actor
- **fragile**: suitable in case minimal alterations of the carrier need to be revealed or tracked

Basic (5)

- Two “pillars” of **information hiding** are(*):
 - the needle in the haystack principle
 - the magic triangle



(*) W. Mazurczyk, L. Caviglione, "Steganography in Modern Smartphones and Mitigation Techniques", IEEE Communications Surveys & Tutorials, Vol. 17, No. 1, pp. 334-357, First Quarter 2015

Examples of Watermarking Techniques

Roadmap

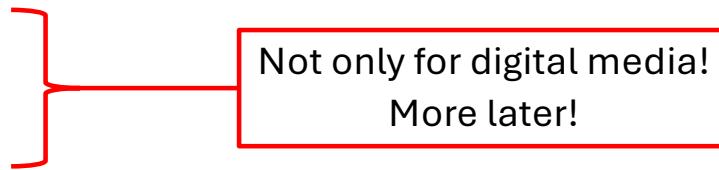
- The various watermarking techniques will be explored via **examples**:
 - digital media
 - 3D mesh
 - network traffic
 - software
 - hardware
 - AI models and outputs
- Choice:
 - showing a mechanism from **different points of view**
 - emphasizing the **hiding approach - carrier** relationship
 - have some nice **pointers** for further investigation
 - have fun!



Digital Media (1)

- Digital media are **ubiquitous**:
 - medical
 - entertainment
 - applications assets, e.g., icons and artworks
 - ...
- The literature offers a **huge** variety of watermarking methods:
 - each one tailored for a specific goal or carrier
- **Two** main categories:
 - spatial domain
 - frequency domain (e.g., DCT and DFT)

Digital Media (2)

- Least Significant Bit (LSB) watermarking is the most **famous** family of methods:
 - easy yet effective
 - very popular and portable
 - Directly rooted from LSB steganography, it can be used with:
 - digital pictures
 - digital video/audio
 - timestamps
 - numerical values
 - ...
- 
- Not only for digital media!
More later!

Digital Media (2)

- LSB watermark requires the carrier to **resist** a single-bit modification:
 - not suitable for software or “structured data”!

```
000000d0 20 20 20 20 20 20 20 0a 32 36 20 30 20 6f 62 6a | .26 0 obj|  
000000e0 0a 3c 3c 20 2f 54 79 70 65 20 2f 58 52 65 66 20 | .<< /Type /XRef |  
000000f0 2f 4c 65 6e 67 74 68 20 37 37 20 2f 46 69 6c 74 | /Length 77 /Filt |  
00000100 65 72 20 2f 46 6c 61 74 65 44 65 63 6f 64 65 20 | /er /FlateDecode |  
00000110 2f 44 65 63 6f 64 65 50 61 72 6d 73 20 3c 3c 20 | /DecodeParms << |  
00000120 2f 43 6f 6c 75 6d 6e 73 20 35 20 2f 50 72 65 64 | /Columns 5 /Pred |  
00000130 69 63 74 6f 72 20 31 32 20 3e 3e 20 2f 57 20 5b | /ictor 12 >> /W [ |  
00000140 20 31 20 33 20 31 20 5d 20 2f 49 6e 64 65 78 20 | 1 3 1 ] /Index |  
00000150 5b 20 32 35 20 35 37 20 5d 20 2f 49 6e 66 6f 20 | [ 25 57 ] /Info |  
00000160 32 33 20 30 20 52 20 2f 52 6f 6f 74 20 32 37 20 | 123 0 R /Root 27 |  
00000170 30 20 52 20 2f 53 69 7a 65 20 38 32 20 2f 50 72 | 10 R /Size 82 /Pr |  
00000180 65 76 20 34 31 35 38 34 39 20 20 20 20 20 20 20 | /ev 415649 |  
00000190 20 20 20 20 20 20 20 20 20 2f 49 44 20 5t 3c 31 | /ID [<1|  
000001a0 35 36 32 33 61 62 65 30 63 39 35 33 62 64 62 39 | 15623abe0c953bdb91 |  
000001b0 65 36 64 63 32 63 30 65 36 34 33 34 66 61 39 3e | e6dc2c0e6434fa9>|  
000001c0 3c 62 30 36 39 33 32 35 36 64 65 36 34 65 38 37 | <b0693256de64e87 |  
000001d0 65 64 35 37 66 66 35 66 35 63 34 61 31 35 65 62 | ed57ff5f5c4a15eb |  
000001e0 38 3e 5d 20 3e 3e 0a 73 74 72 65 61 6d 0a 78 9c | 18>] >>.stream.x. |  
000001f0 63 62 64 60 e0 67 60 62 60 60 38 09 22 99 cf 81 | lcbd` .g` b` `8. .... |  
00000200 d9 c6 20 92 f1 28 98 cc 03 91 42 45 20 52 6e 3b | ... ...C....BE Rn; |  
00000210 58 0d 2b 88 e4 ca 01 91 86 2a 20 d2 bc 03 44 7a | IX.+.....* ...Dz |  
00000220 ff 06 92 8c 16 0c 60 35 2c a8 24 e3 ff 53 27 c0 | .....`5,.$.S'. |  
00000230 26 33 30 8e 04 12 00 85 e2 0a bd 0a 65 6e 64 73 | 1830.....ends |  
00000240 74 72 65 61 6d 0a 65 6e 64 6f 62 6a 0a 20 20 20 | Itream.endobj. |  
00000250 20 20 20 20 20 20 20 20 20 20 20 20 20 20 20 20 | |
```

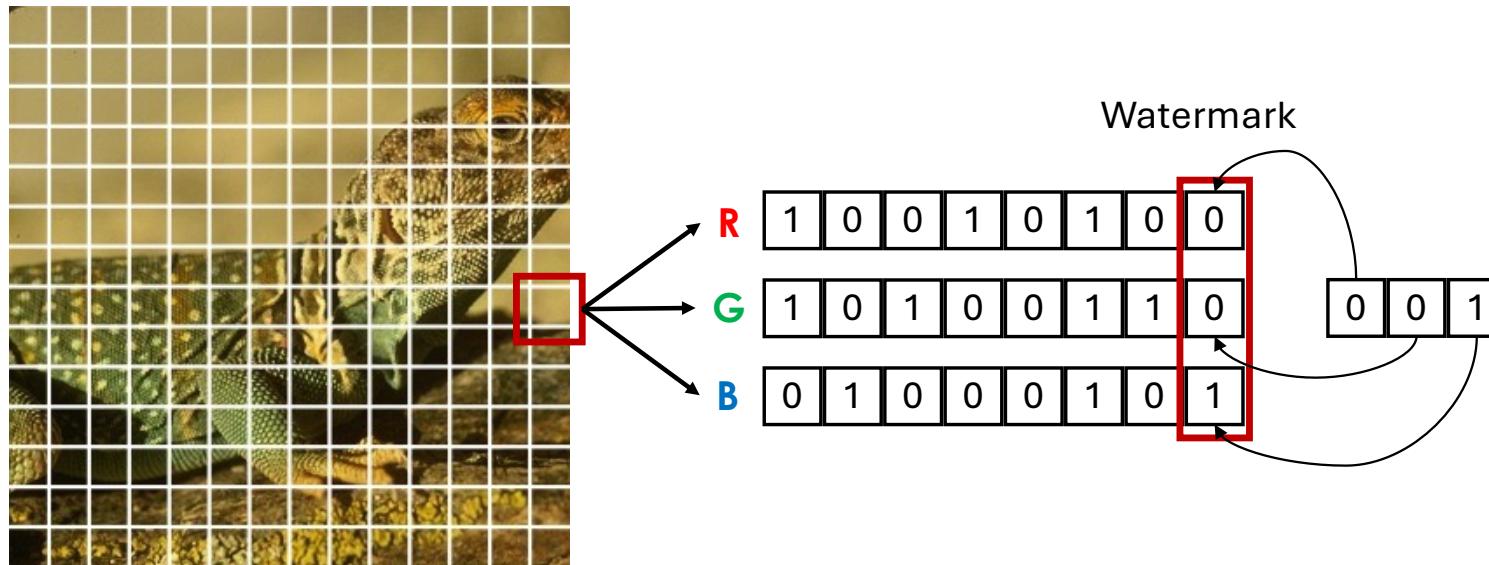
Hex Dump of .pdf

Altering a bit will mostly render the document unreadable
(not to mention that it is easy to spot by computing the MD5)

LSB Watermarking (1)

- As a paradigmatic example, will consider digital pictures

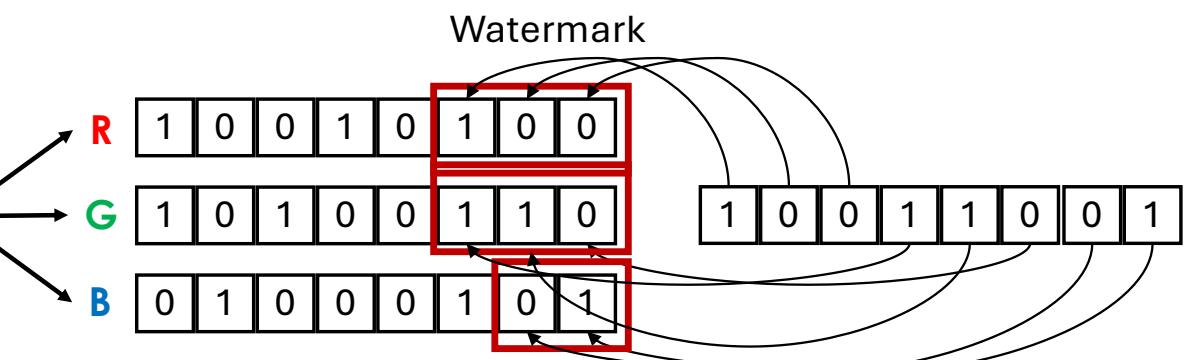
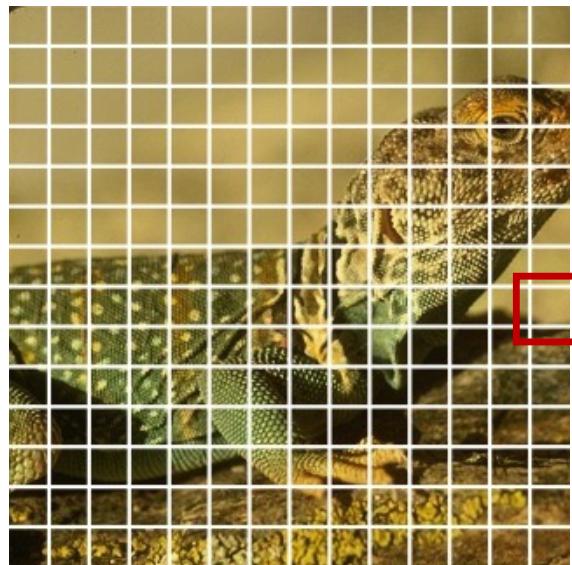
LSB “standard”



LSB Watermarking (1)

- As a paradigmatic example, will consider digital pictures

LSB Multi-Bit



LSB Watermarking (2)

- Demo (1 Bit – Red Channel):



Original Image



Watermarked Image

Secret: 100 times "Summer School on Artificial Intelligence for a Secure Society"

LSB Watermarking (2)

- Demo (1 Bit – Red Channel):

Blue Channel
(Grayscale)



Original Image



Watermarked Image

LSB Watermarking (2)

- Demo (1 Bit – Red Channel):

Green Channel
(Grayscale)



Original Image



Watermarked Image

LSB Watermarking (2)

- Demo (1 Bit – Red Channel):

Red Channel
(Grayscale)



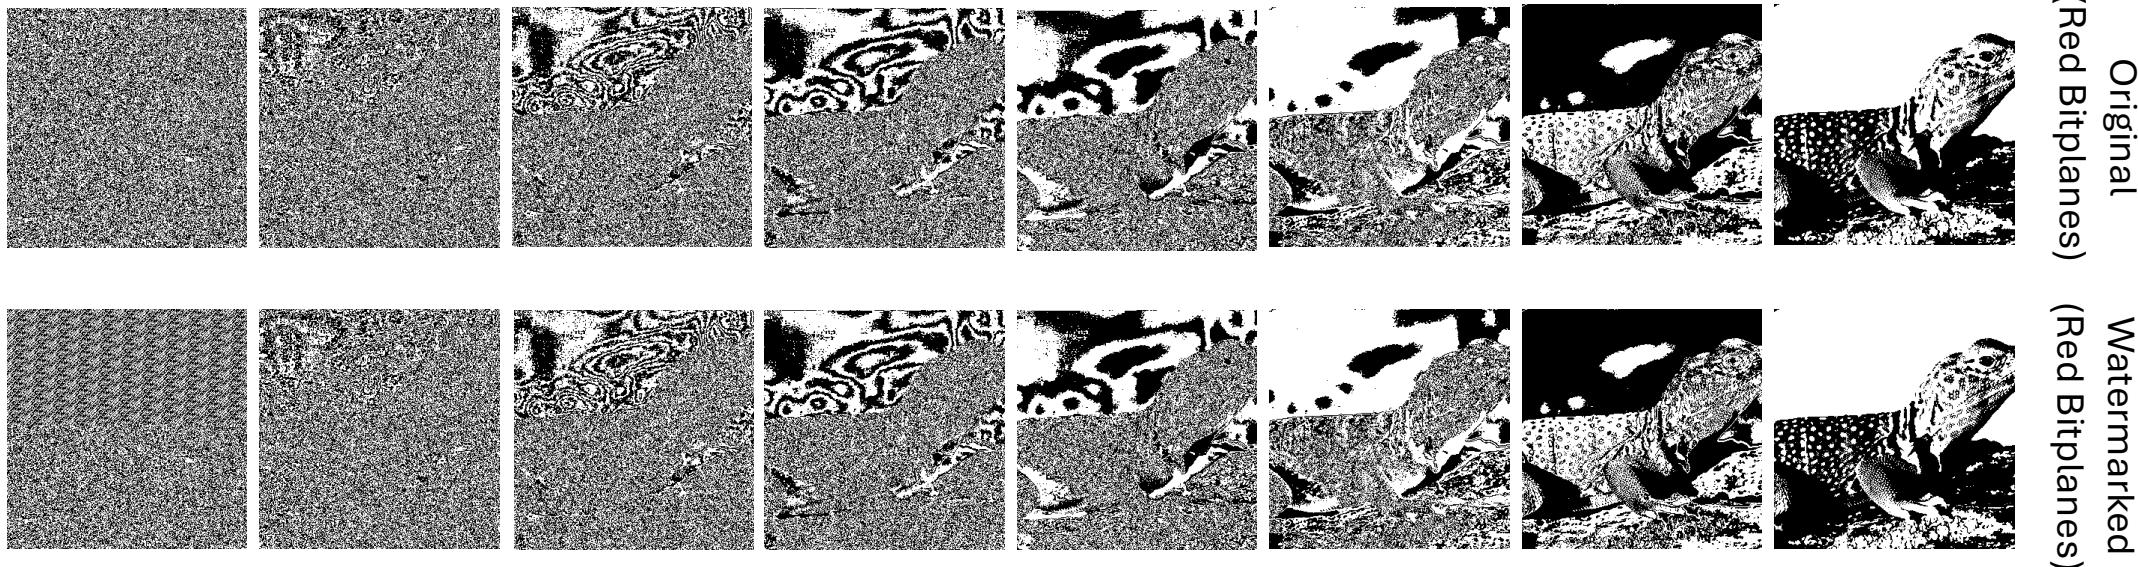
Original Image



Watermarked Image

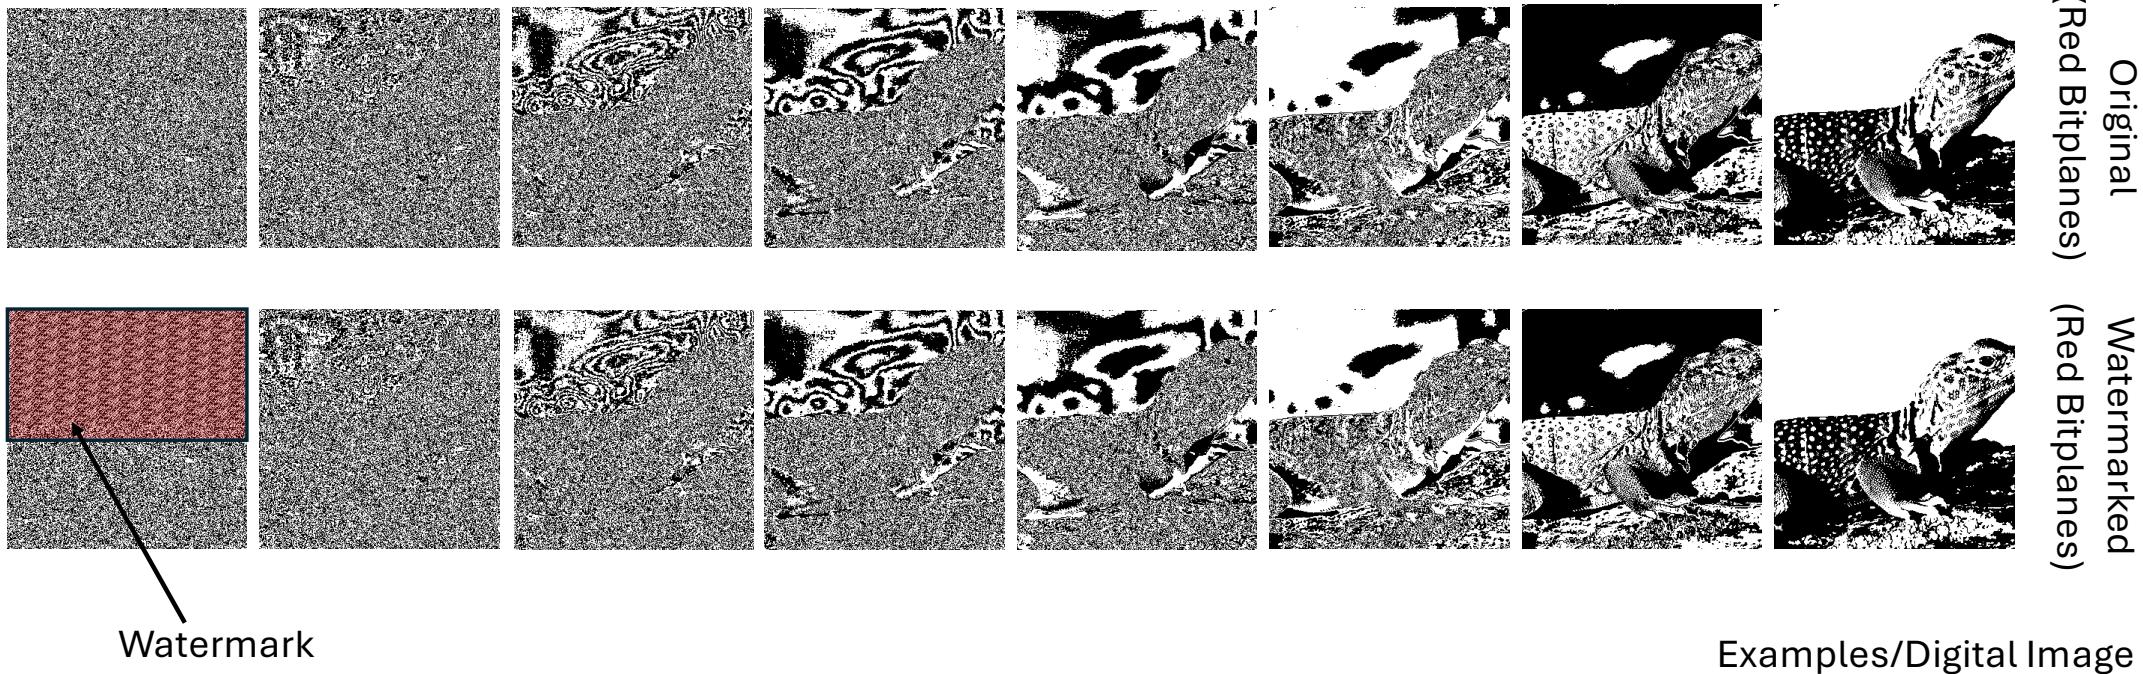
LSB Watermarking (2)

- Demo (1 Bit – Red Channel):



LSB Watermarking (2)

- Demo (1 Bit – Red Channel):



3D Mesh (1)

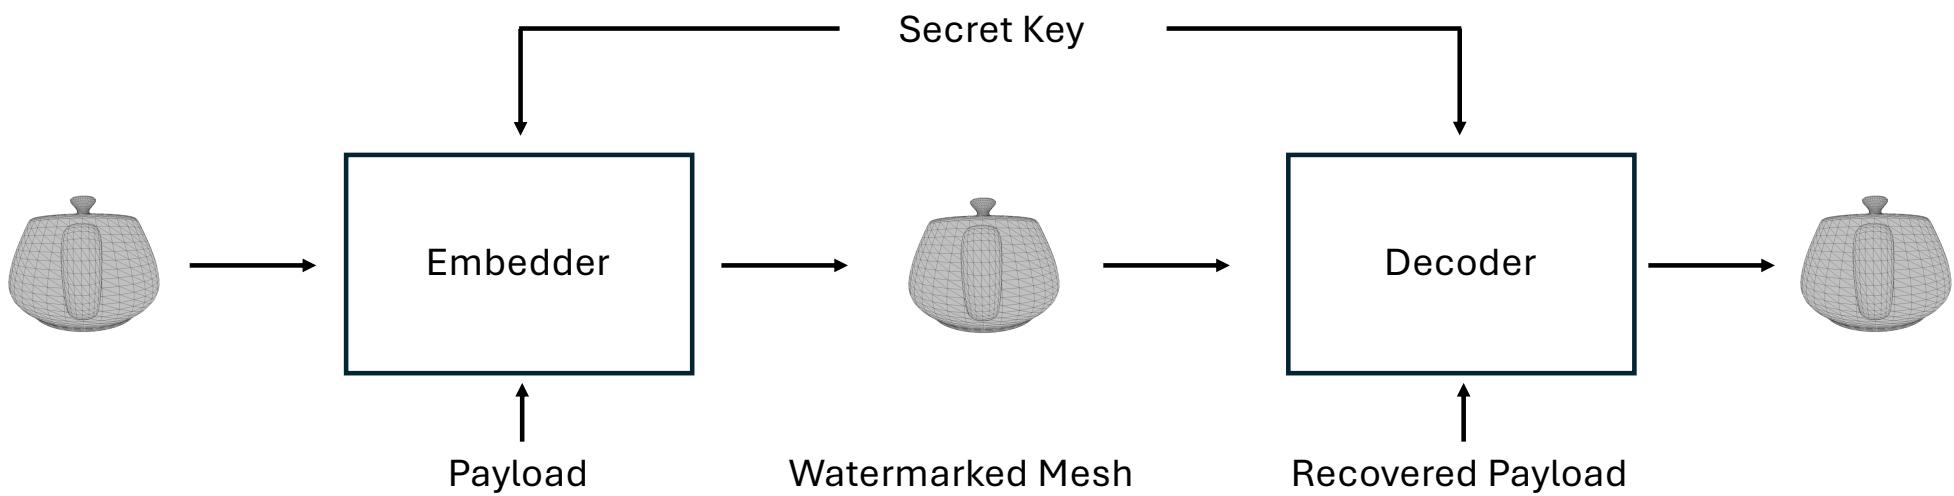
- 3D meshes are at the **basis** of many emerging applications:
 - virtual reality
 - video games / entertainment
 - medical imaging
 - CAD
 - printing
- A **mesh** is composed of:
 - vertices, edges, and facets
 - geometry (3D positions via coordinates)
 - connectivity (adjacency relations among elements)
- More complex data requires a more complex approach

3D Mesh (2)

- The most valuable meshes are:
 - **difficult** to obtain (e.g., require the work of several artists)
 - very **valuable** (obtained through expensive data gathering/scanning campaigns)
- Literature proposes various watermarking techniques
- **Fragile** methods:
 - typically used for authentication, should disrupt against minimal modifications
 - should locate or identify the attack
 - can hide data in spatial (connectivity) or transformed (spectral) domains
- **Robust** methods:
 - typically used for IP management, should resist against great distortions/degradations
 - should balance robustness and secrecy
 - can hide data in spatial (connectivity) or transformed (spectral) domains

3D Triangle Mesh (1)

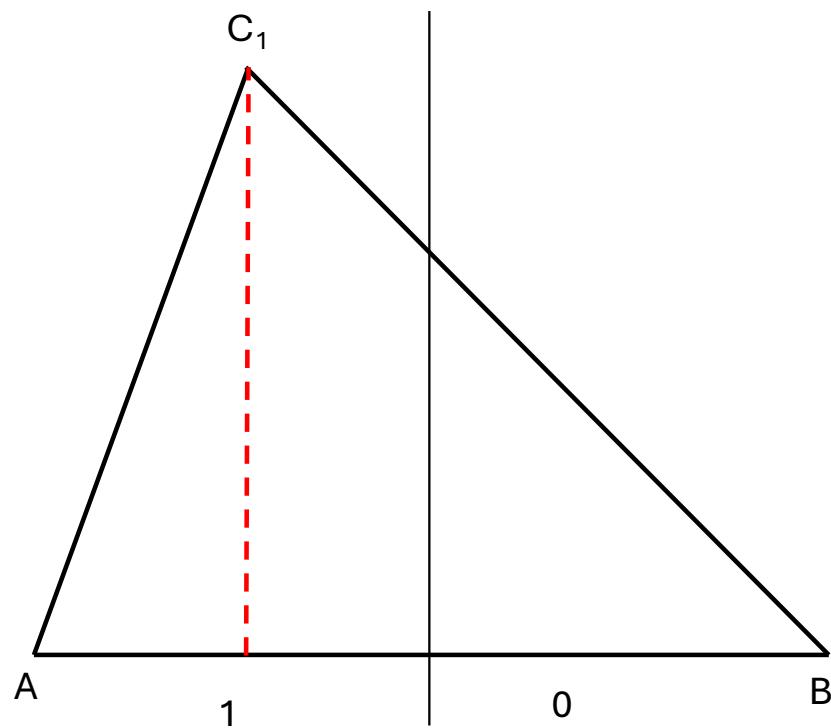
- Example of a **fragile** technique in the **spatial** domain:
 - based on the modification of the geometry
 - for triangle meshes
 - method by Cayre and Macq (*)



(*) F. Cayre, B. Macq, "Data hiding on 3-D triangle meshes", IEEE Transactions on Signal Processing, Vol. 51, No. 4, pp. 939-949, April 2003

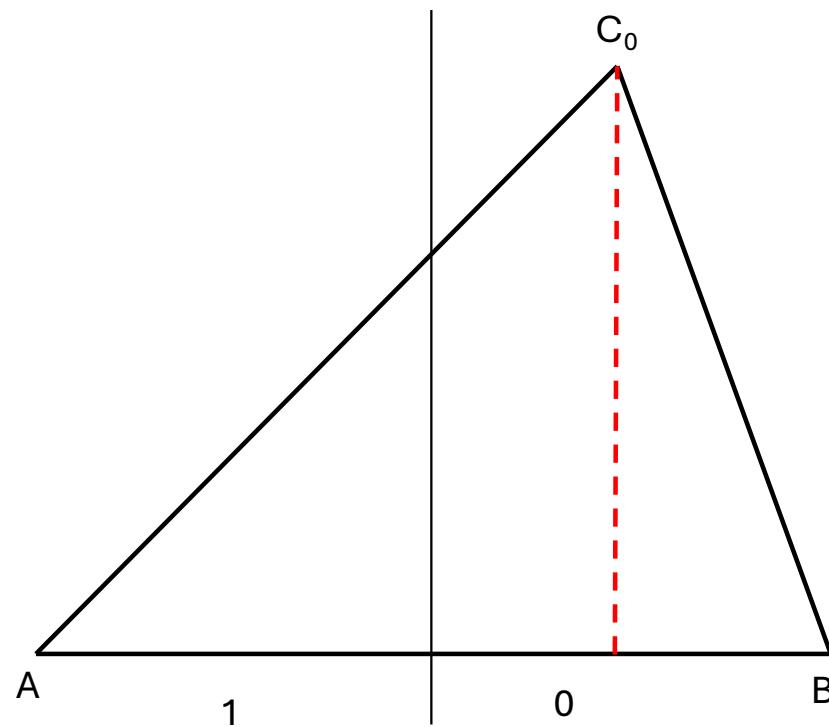
3D Triangle Mesh (2)

- Key idea:



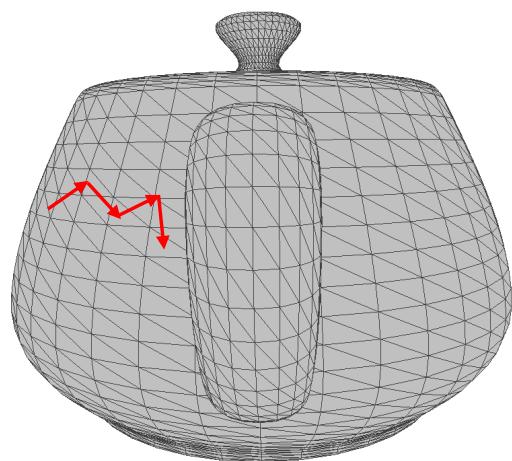
3D Triangle Mesh (2)

- Key idea:

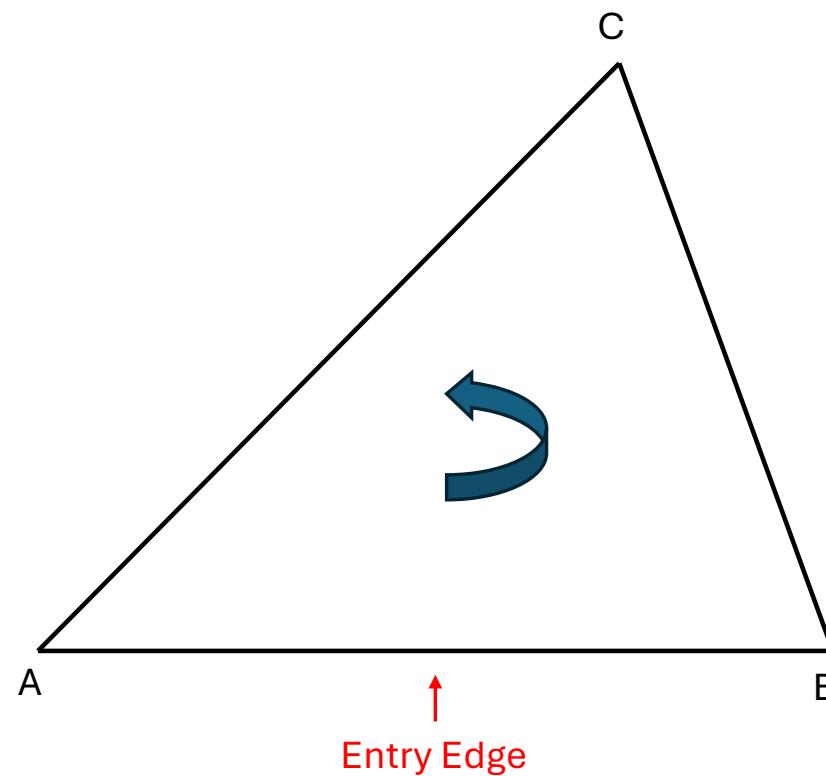


3D Triangle Mesh (2)

- How to “navigate” the mesh:

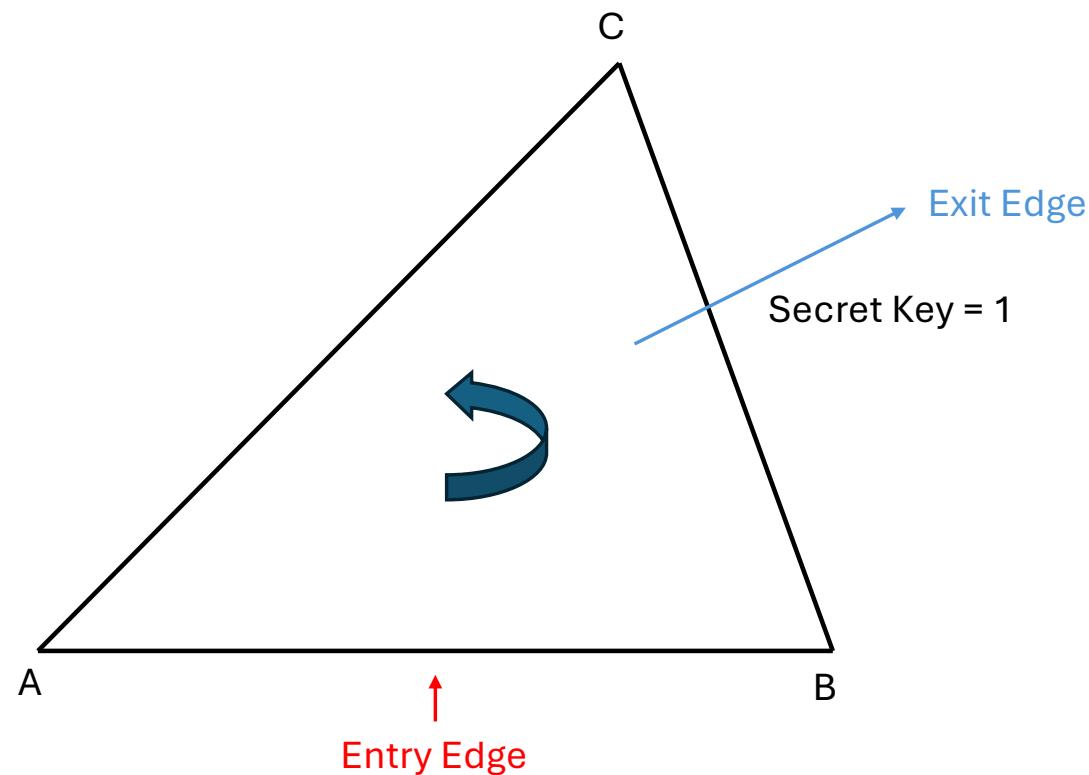
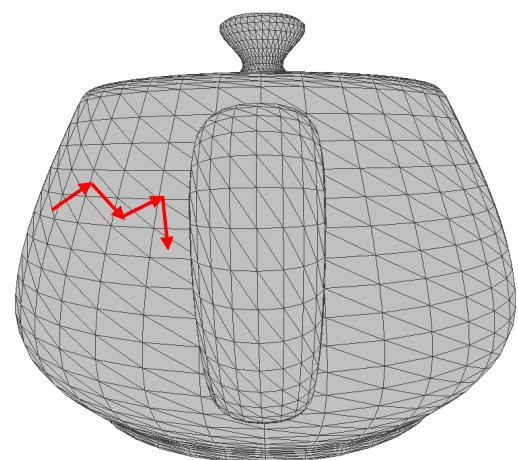


The starting triangle is selected according to geometric constraints



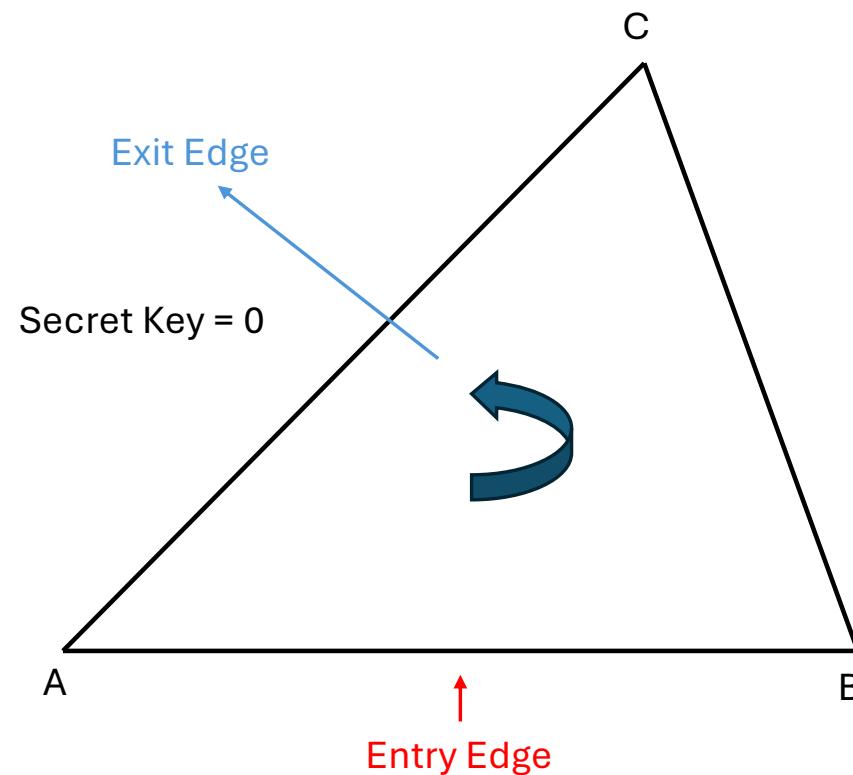
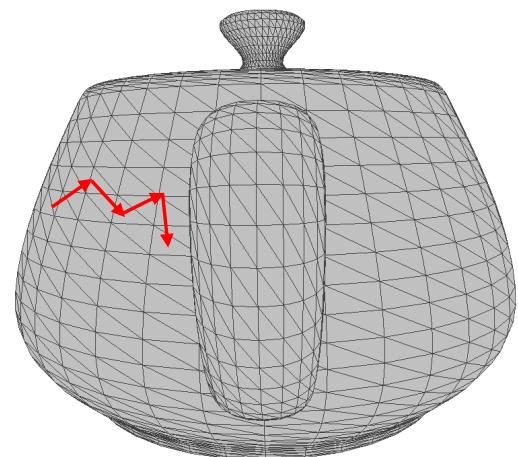
3D Triangle Mesh (2)

- How to “navigate” the mesh:



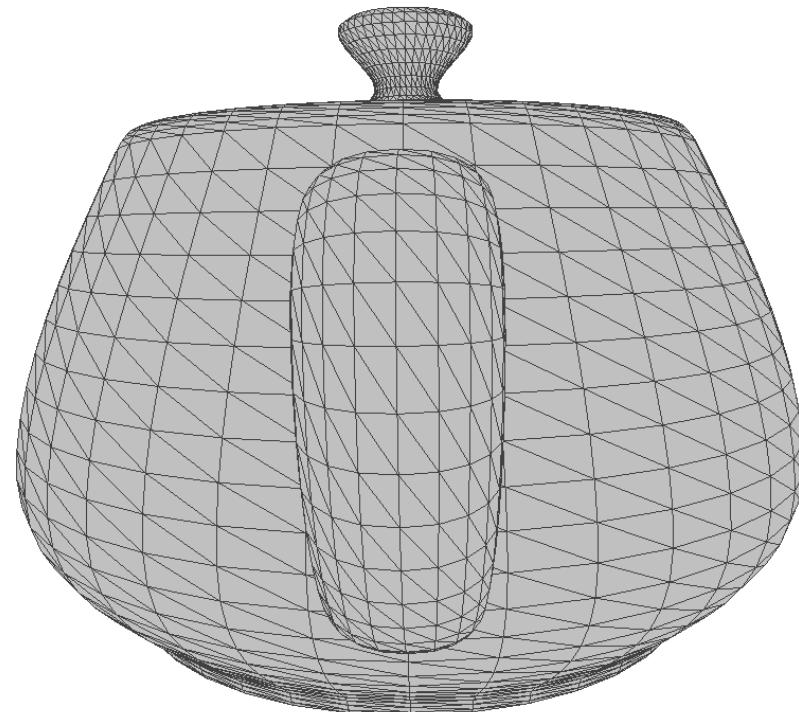
3D Triangle Mesh (2)

- How to “navigate” the mesh:



3D Triangle Mesh (3)

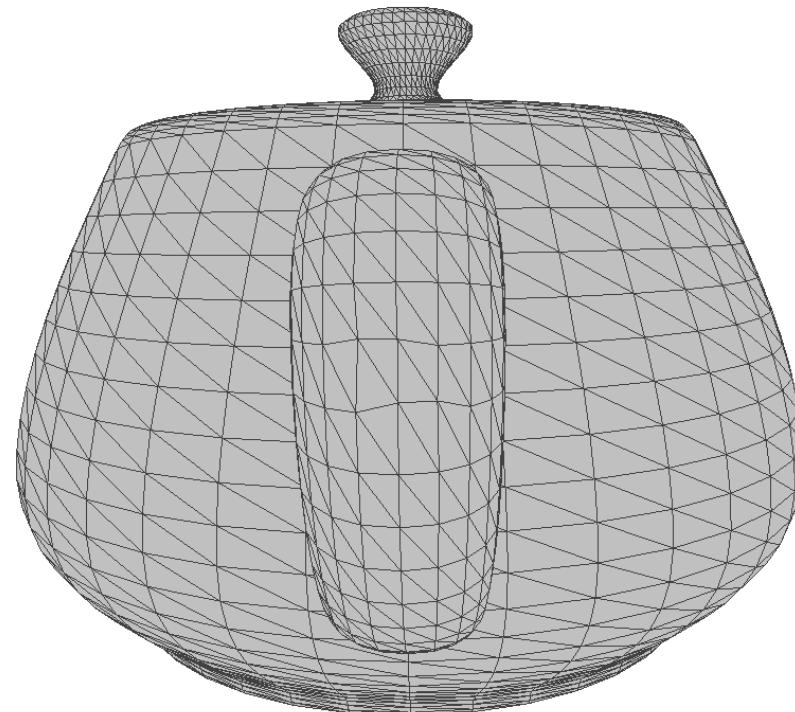
- Demo:



Examples/3D Triangle Mesh

3D Triangle Mesh (3)

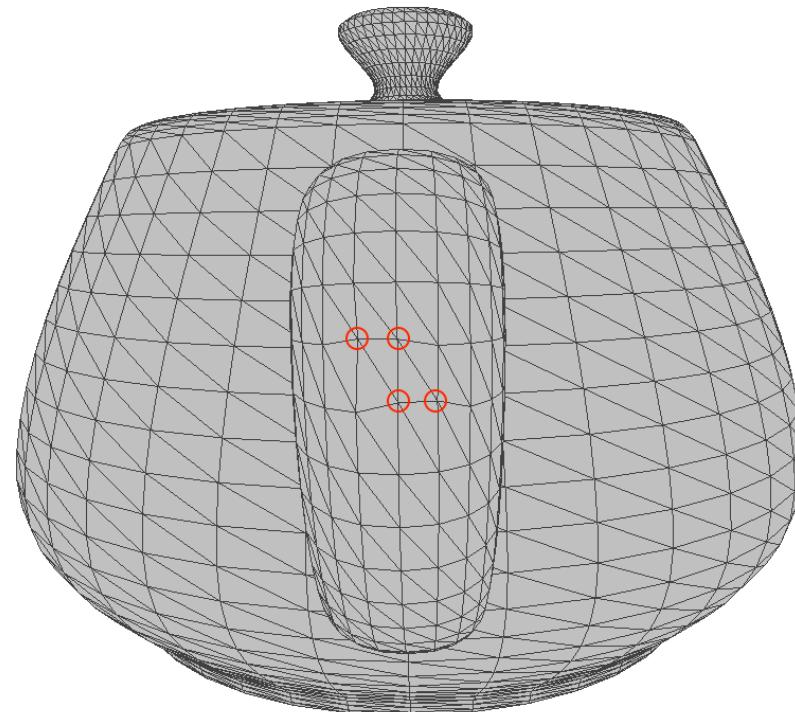
- Demo:



Examples/3D Triangle Mesh

3D Triangle Mesh (3)

- Demo:



Examples/3D Triangle Mesh

Network Traffic (1)

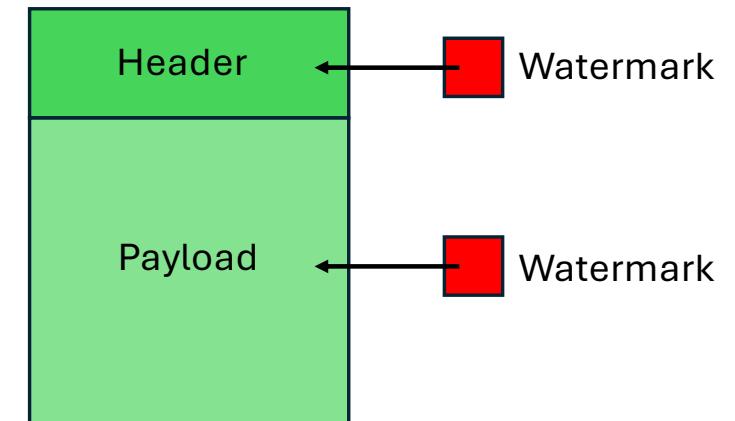
- Network traffic is **ubiquitous** as the Internet:
 - “glue” for microservices
 - contains sensitive data (even if encrypted)
 - re-routed by threat actors
 - processed by many defensive tools (even AI based)
- There is the need of techniques for:
 - traffic engineering
 - network design / troubleshoot
 - tracing malware, attackers, and botnets
 - identify endpoints even when encryption is used

Network Traffic (2)

- Compared to other domains, **traffic watermarking** is quite **new**
- A good number of mechanisms is becoming available(*):
 - techniques are often borrowed from the literature on **covert channels**
- They are usually classified according to **where** the watermark is hidden

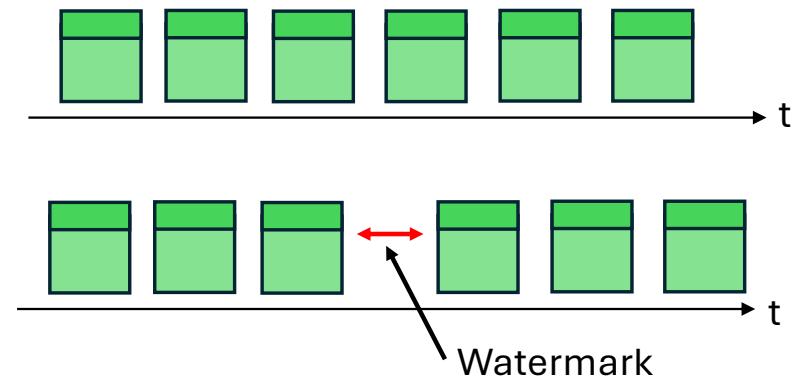
Network Traffic (2)

- Compared to other domains, **traffic watermarking** is quite **new**
- A good number of mechanisms is becoming available:
 - techniques are often borrowed from the literature on **covert channels**
- They are usually classified according to **where** the watermark is hidden
- **Content** based:
 - header
 - payload



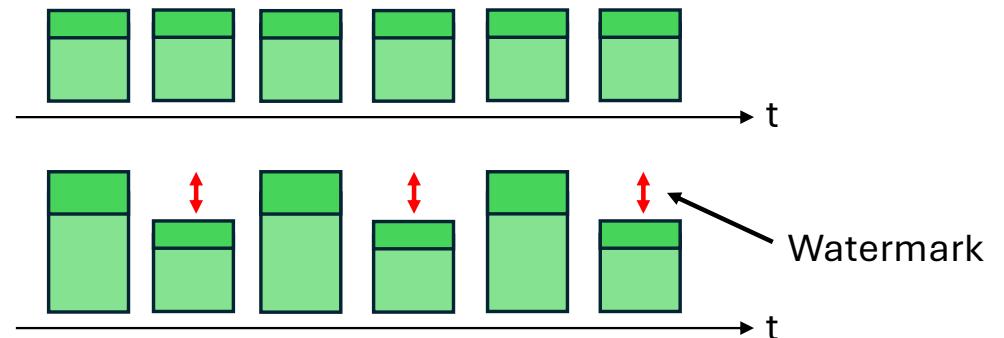
Network Traffic (2)

- Compared to other domains, **traffic watermarking** is quite **new**
- A good number of mechanisms is becoming available:
 - techniques are often borrowed from the literature on **covert channels**
- They are usually classified according to **where** the watermark is hidden
- **Timing** based:
 - delay
 - counting
 - statistical



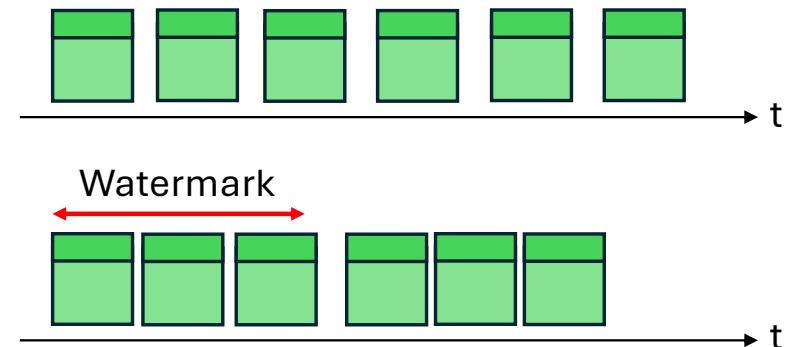
Network Traffic (2)

- Compared to other domains, **traffic watermarking** is quite **new**
- A good number of mechanisms is becoming available:
 - techniques are often borrowed from the literature on **covert channels**
- They are usually classified according to **where** the watermark is hidden
- **Size** based:
 - packet length
 - optional fields



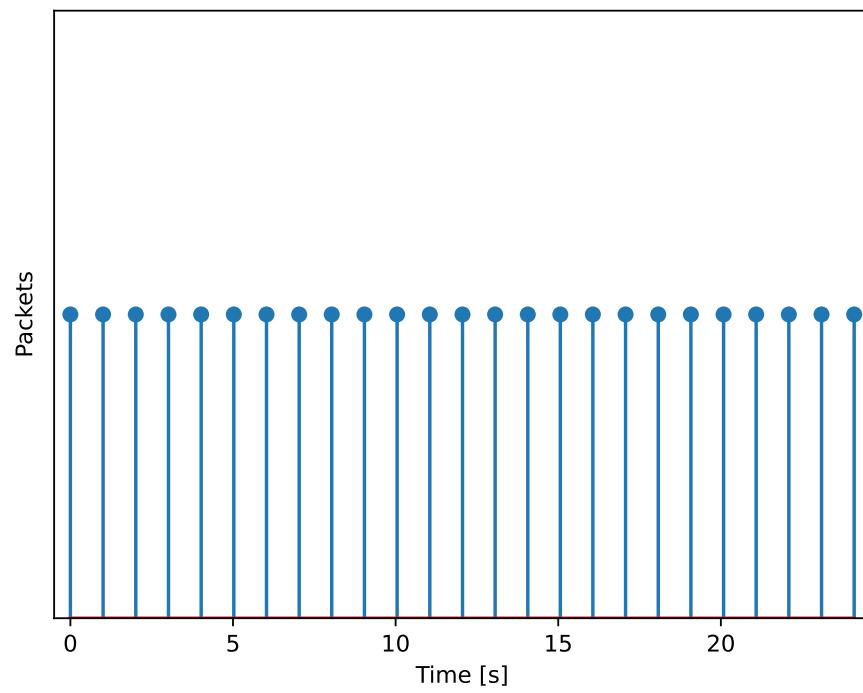
Network Traffic (2)

- Compared to other domains, **traffic watermarking** is quite **new**
- A good number of mechanisms is becoming available:
 - techniques are often borrowed from the literature on **covert channels**
- They are usually classified according to **where** the watermark is hidden
- **Rate** based:
 - throughput
 - packet rate

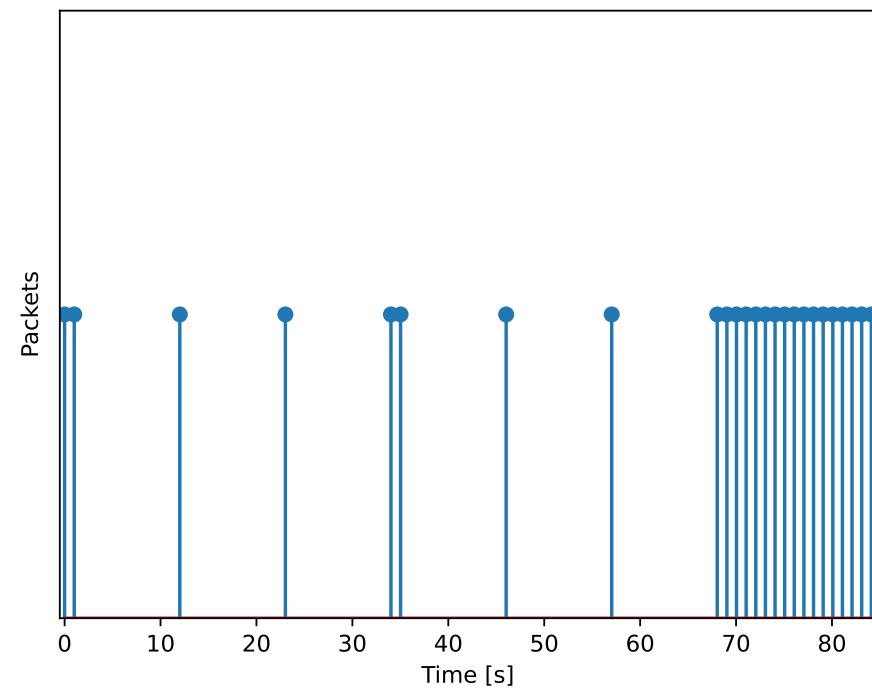


Network Traffic (3)

- Demo (timing watermark):



Inter-Packet Time of the Original Traffic

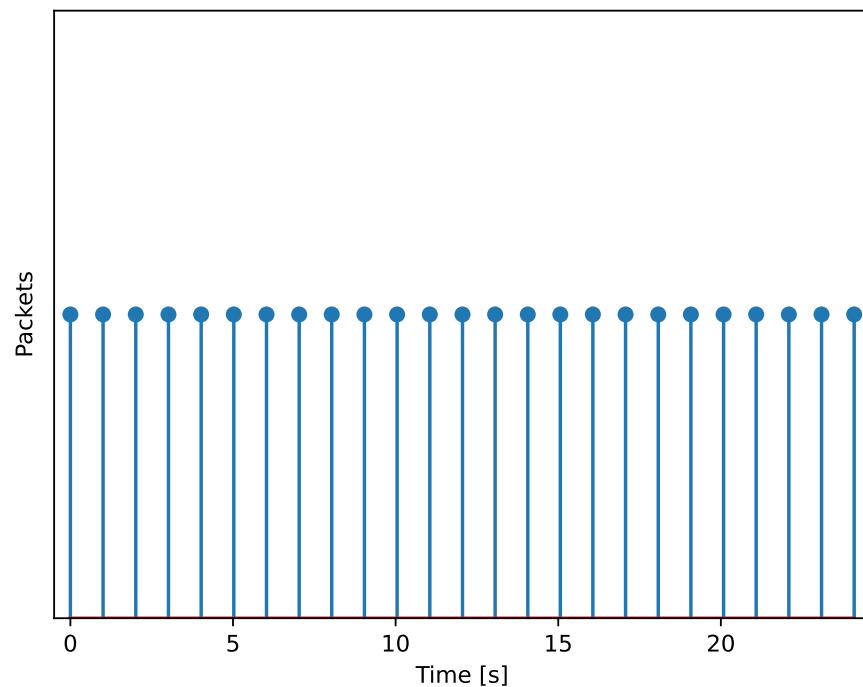


Inter-Packet Time of the Watermarked Traffic

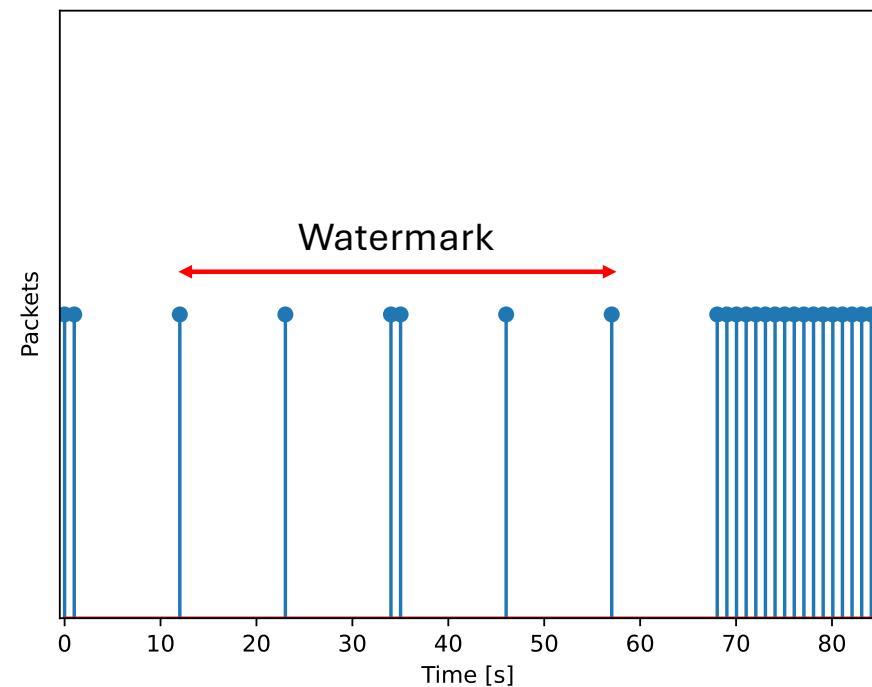
Examples/Network Traffic/Timing Method

Network Traffic (3)

- Demo (timing watermark):



Inter-Packet Time of the Original Traffic



Inter-Packet Time of the Watermarked Traffic

Examples/Network Traffic/Timing Method

Network Traffic (3)

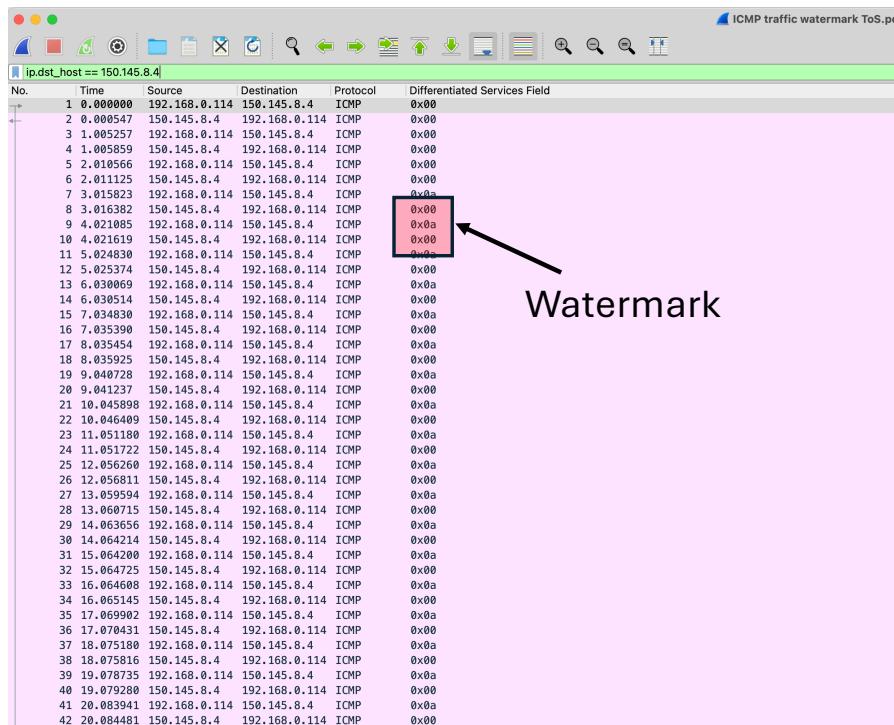
- Demo (content watermark):

No.	Time	Source	Destination	Protocol	Differentiated Services Field
1	0.000000	192.168.0.114	150.145.8.4	ICMP	0x00
2	0.000547	150.145.8.4	192.168.0.114	ICMP	0x00
3	1.005257	192.168.0.114	150.145.8.4	ICMP	0x00
4	1.005859	150.145.8.4	192.168.0.114	ICMP	0x00
5	2.010566	192.168.0.114	150.145.8.4	ICMP	0x00
6	2.011125	150.145.8.4	192.168.0.114	ICMP	0x00
7	3.015823	192.168.0.114	150.145.8.4	ICMP	0x0a
8	3.016382	150.145.8.4	192.168.0.114	ICMP	0x00
9	4.021085	192.168.0.114	150.145.8.4	ICMP	0x0a
10	4.021619	150.145.8.4	192.168.0.114	ICMP	0x00
11	5.024830	192.168.0.114	150.145.8.4	ICMP	0x0a
12	5.025374	150.145.8.4	192.168.0.114	ICMP	0x00
13	6.030063	192.168.0.114	150.145.8.4	ICMP	0x0a
14	6.030514	150.145.8.4	192.168.0.114	ICMP	0x00
15	7.034839	192.168.0.114	150.145.8.4	ICMP	0x0a
16	7.035390	150.145.8.4	192.168.0.114	ICMP	0x00
17	8.035454	192.168.0.114	150.145.8.4	ICMP	0x0a
18	8.035925	150.145.8.4	192.168.0.114	ICMP	0x00
19	9.040728	192.168.0.114	150.145.8.4	ICMP	0x0a
20	9.041237	150.145.8.4	192.168.0.114	ICMP	0x00
21	10.045898	192.168.0.114	150.145.8.4	ICMP	0x0a
22	10.046409	150.145.8.4	192.168.0.114	ICMP	0x00
23	11.051188	192.168.0.114	150.145.8.4	ICMP	0x0a
24	11.051722	150.145.8.4	192.168.0.114	ICMP	0x00
25	12.056260	192.168.0.114	150.145.8.4	ICMP	0x0a
26	12.056811	150.145.8.4	192.168.0.114	ICMP	0x00
27	13.059594	192.168.0.114	150.145.8.4	ICMP	0x0a
28	13.060715	150.145.8.4	192.168.0.114	ICMP	0x00
29	14.063656	192.168.0.114	150.145.8.4	ICMP	0x0a
30	14.064214	150.145.8.4	192.168.0.114	ICMP	0x00
31	15.064200	192.168.0.114	150.145.8.4	ICMP	0x0a
32	15.064725	150.145.8.4	192.168.0.114	ICMP	0x00
33	16.064608	192.168.0.114	150.145.8.4	ICMP	0x0a
34	16.065145	150.145.8.4	192.168.0.114	ICMP	0x00
35	17.069902	192.168.0.114	150.145.8.4	ICMP	0x0a
36	17.070431	150.145.8.4	192.168.0.114	ICMP	0x00
37	18.075100	192.168.0.114	150.145.8.4	ICMP	0x0a
38	18.075816	150.145.8.4	192.168.0.114	ICMP	0x00
39	19.078735	192.168.0.114	150.145.8.4	ICMP	0x0a
40	19.079280	150.145.8.4	192.168.0.114	ICMP	0x00
41	20.083941	192.168.0.114	150.145.8.4	ICMP	0x0a
42	20.084481	150.145.8.4	192.168.0.114	ICMP	0x00

Examples/Network Traffic/Storage Methods

Network Traffic (3)

- Demo (content watermark):



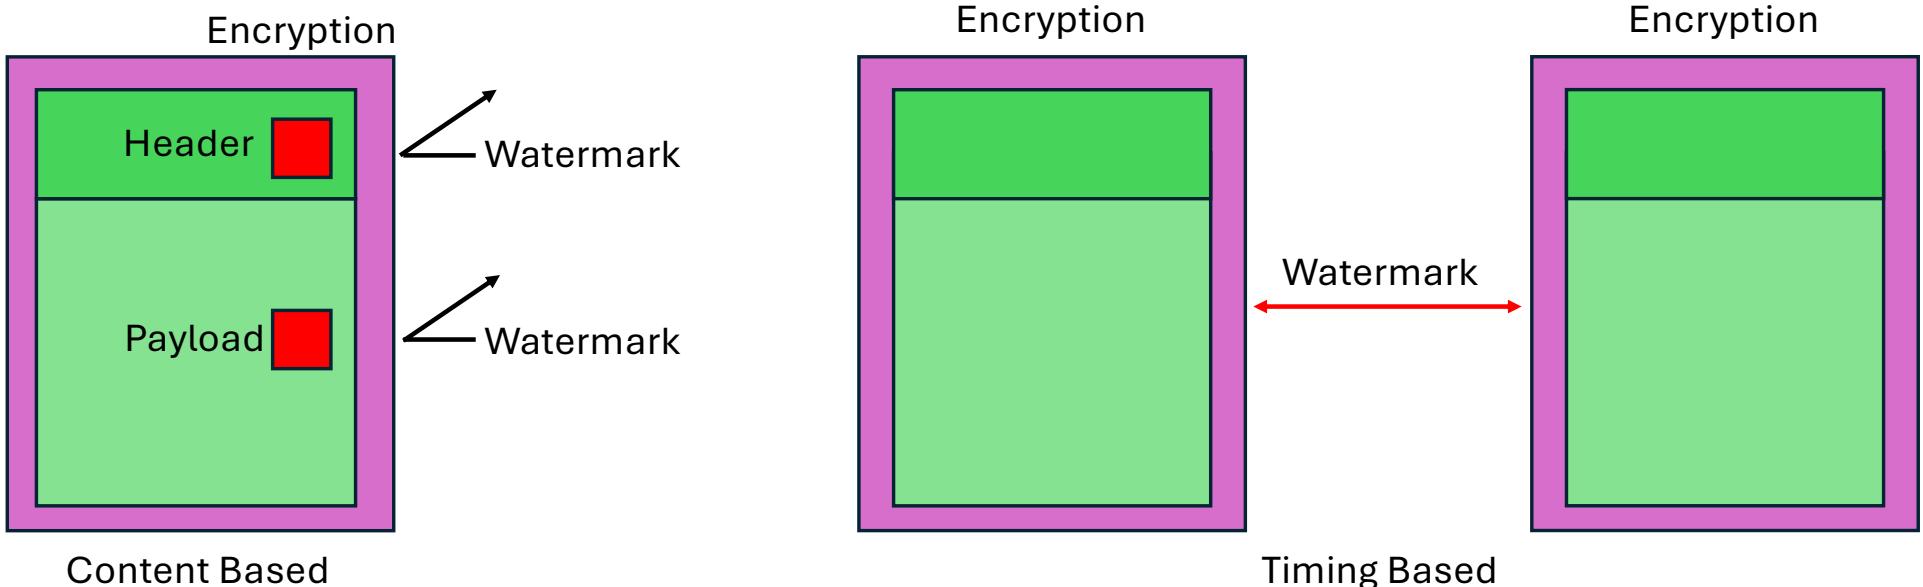
Watermark

No.	Time	Source	Destination	Protocol	Differentiated Services Field
1	0.000000	192.168.0.114	150.145.8.4	ICMP	0x00
2	0.000547	150.145.8.4	192.168.0.114	ICMP	0x00
3	1.005257	192.168.0.114	150.145.8.4	ICMP	0x00
4	1.005859	150.145.8.4	192.168.0.114	ICMP	0x00
5	2.010566	192.168.0.114	150.145.8.4	ICMP	0x00
6	2.011125	150.145.8.4	192.168.0.114	ICMP	0x00
7	3.015823	192.168.0.114	150.145.8.4	ICMP	0x00
8	3.016382	150.145.8.4	192.168.0.114	ICMP	0x0a
9	4.021085	192.168.0.114	150.145.8.4	ICMP	0x00
10	4.021619	150.145.8.4	192.168.0.114	ICMP	0x00
11	5.024830	192.168.0.114	150.145.8.4	ICMP	0x00
12	5.025374	150.145.8.4	192.168.0.114	ICMP	0x00
13	6.030063	192.168.0.114	150.145.8.4	ICMP	0x0a
14	6.030514	150.145.8.4	192.168.0.114	ICMP	0x00
15	7.034839	192.168.0.114	150.145.8.4	ICMP	0x0a
16	7.035390	150.145.8.4	192.168.0.114	ICMP	0x00
17	8.035454	192.168.0.114	150.145.8.4	ICMP	0x0a
18	8.035925	150.145.8.4	192.168.0.114	ICMP	0x00
19	9.040728	192.168.0.114	150.145.8.4	ICMP	0x0a
20	9.041237	150.145.8.4	192.168.0.114	ICMP	0x00
21	10.045898	192.168.0.114	150.145.8.4	ICMP	0x0a
22	10.046409	150.145.8.4	192.168.0.114	ICMP	0x00
23	11.051188	192.168.0.114	150.145.8.4	ICMP	0x0a
24	11.051722	150.145.8.4	192.168.0.114	ICMP	0x00
25	12.056260	192.168.0.114	150.145.8.4	ICMP	0x0a
26	12.056811	150.145.8.4	192.168.0.114	ICMP	0x00
27	13.059594	192.168.0.114	150.145.8.4	ICMP	0x0a
28	13.060715	150.145.8.4	192.168.0.114	ICMP	0x00
29	14.063656	192.168.0.114	150.145.8.4	ICMP	0x0a
30	14.064214	150.145.8.4	192.168.0.114	ICMP	0x00
31	15.064200	192.168.0.114	150.145.8.4	ICMP	0x0a
32	15.064725	150.145.8.4	192.168.0.114	ICMP	0x00
33	16.064608	192.168.0.114	150.145.8.4	ICMP	0x0a
34	16.065145	150.145.8.4	192.168.0.114	ICMP	0x00
35	17.069902	192.168.0.114	150.145.8.4	ICMP	0x0a
36	17.070431	150.145.8.4	192.168.0.114	ICMP	0x00
37	18.075180	192.168.0.114	150.145.8.4	ICMP	0x0a
38	18.075816	150.145.8.4	192.168.0.114	ICMP	0x00
39	19.078735	192.168.0.114	150.145.8.4	ICMP	0x0a
40	19.079280	150.145.8.4	192.168.0.114	ICMP	0x00
41	20.083941	192.168.0.114	150.145.8.4	ICMP	0x0a
42	20.084481	150.145.8.4	192.168.0.114	ICMP	0x00

Examples/Network Traffic/Storage Methods

Network Traffic (4)

- Each method has:
 - **pros and cons**
 - suitability for a specific **goal**
- Example:



Software (1)

- **Software** is today a **critical** asset (obviously):
 - piracy
 - reverse engineering
 - IP infringements
- With the advent of **AI-based code generators or assistants**:
 - new problems due to the ingestion of source code!
 - ability of generating malicious/unaudited code is magnified!
- Two main mechanisms for marking the software:
 - **static**: the watermark is hidden in code or assets (e.g., icons)
 - **dynamic**: the watermark is stored in the execution state of a program



Software (2)

- **Three** main exemplary sets of techniques
- **Easter Egg:**
 - simple and “catchy”
 - easy to locate
 - example: *about:mozilla*



Dan Forden

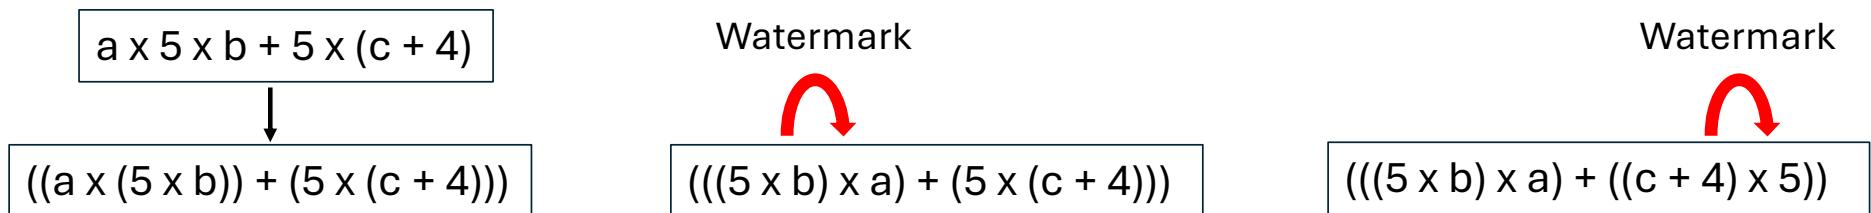
Software (2)

- **Three** main exemplary sets of techniques
- Easter Egg:
 - simple and “catchy”
 - easy to locate
 - example: *about:mozilla*
- **Code substitution** (or replacement):
 - general idea: substitute a portion of the code with a watermarked value
 - simple example: re-ordering of mathematical operations(*)

(*) J. Hamilton, S. Danicic, “A Survey of Static Software Watermarking”, World Congress on Internet Security, London, UK, pp. 100-107, Feb. 2011

Software (2)

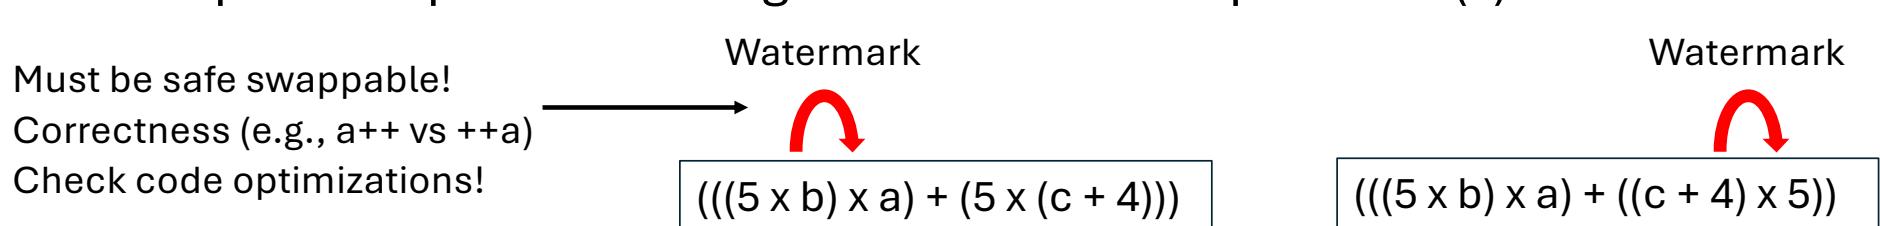
- **Three** main exemplary sets of techniques
- Easter Egg:
 - simple and “catchy”
 - easy to locate
 - example: *about:mozilla*
- **Code substitution** (or replacement):
 - general idea: substitute a portion of the code with a watermarked value
 - simple example: re-ordering of mathematical operations(*)



(*) J. Hamilton, S. Danicic, “A Survey of Static Software Watermarking”, World Congress on Internet Security, London, UK, pp. 100-107, Feb. 2011

Software (2)

- **Three** main exemplary sets of techniques
- Easter Egg:
 - simple and “catchy”
 - easy to locate
 - example: *about:mozilla*
- **Code substitution** (or replacement):
 - general idea: substitute a portion of the code with a watermarked value
 - simple example: re-ordering of mathematical operations(*)



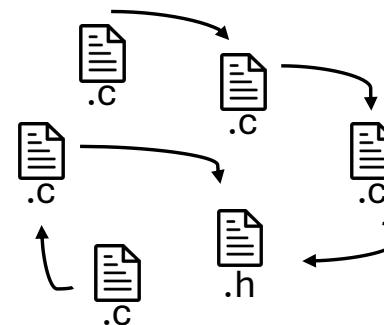
(*) J. Hamilton, S. Danicic, “A Survey of Static Software Watermarking”, World Congress on Internet Security, London, UK, pp. 100-107, Feb. 2011

Software (3)

- **Three** main exemplary sets of techniques
- Easter Egg:
 - simple and “catchy”
 - easy to locate
 - example: *about:mozilla*
- Code substitution (or replacement):
 - general idea: substitute a portion of the code with a watermarked value
 - simple example: re-ordering of mathematical operations(*)
- **Opaque predicates:**
 - borrowed from software protections and obfuscation

Opaque Predicates (1)

- Key idea:
 - create an expression that evaluates to either “**true**” or “**false**”
 - the outcome is **known** a-priori by the **programmer**
 - needs to be evaluated at run time
 - **difficult** to **understand** the intent of the code



Opaque Predicates (1)

- Key idea:
 - create an expression that evaluates to either “true” or “false”
 - the outcome is **known** a-priori by the **programmer**
 - needs to be evaluated at run time
 - **difficult** to **understand** the intent of the code

The screenshot shows a dark-themed code editor with a C program. The code defines variables a, b, and c, calculates d, and then prints "Hello Watermark!" if d equals 8. A red arrow points from the text "Opaque predicate (d is always 8)!" to the if statement. Another red arrow points from the text "Watermark" to the printf call. The code is as follows:

```
1 #include <stdio.h>
2
3 int main()
4 {
5     int a = 3;
6     int b = 10;
7     int c = 34;
8     int d = ((a + b) * c * 2 + 4)/111;
9
10    if (d == 8)
11    {
12        printf("Hello Watermark! \n");
13    }
14
15    return 0;
16 }
```

Opaque predicate
(d is always 8)!

Watermark

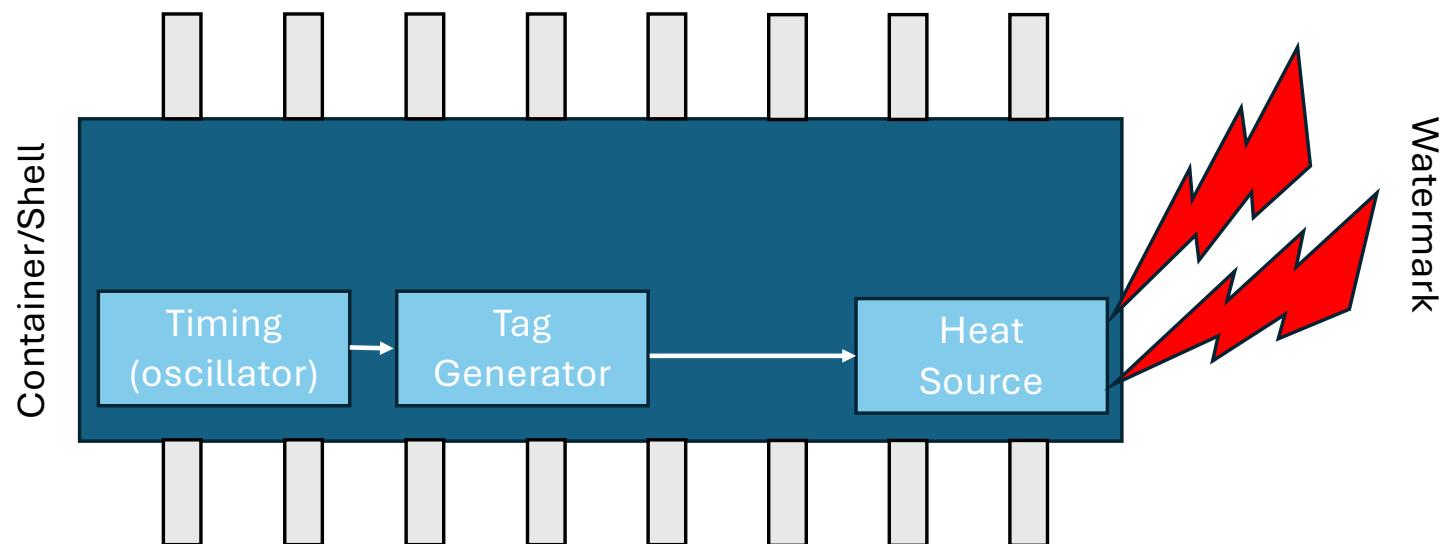
Examples/Opaque Predicate

Hardware (1)

- Custom hardware and “silicon” are now at the basis of:
 - IoT nodes
 - mobile/personal devices
 - CPU/GPU
 - AI-oriented architectures
- Progressive **co-design** of hardware and software for:
 - performances
 - sovereignty
 - greening
- **Good to hear it once :)**

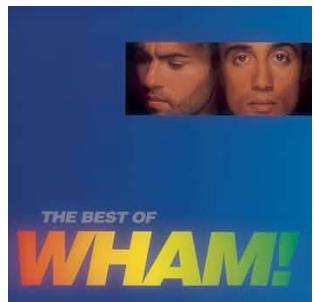
Hardware (2)

- Only mention here a **simple** example:
 - a good starting point is (*) and the references therein



AI Models (1)

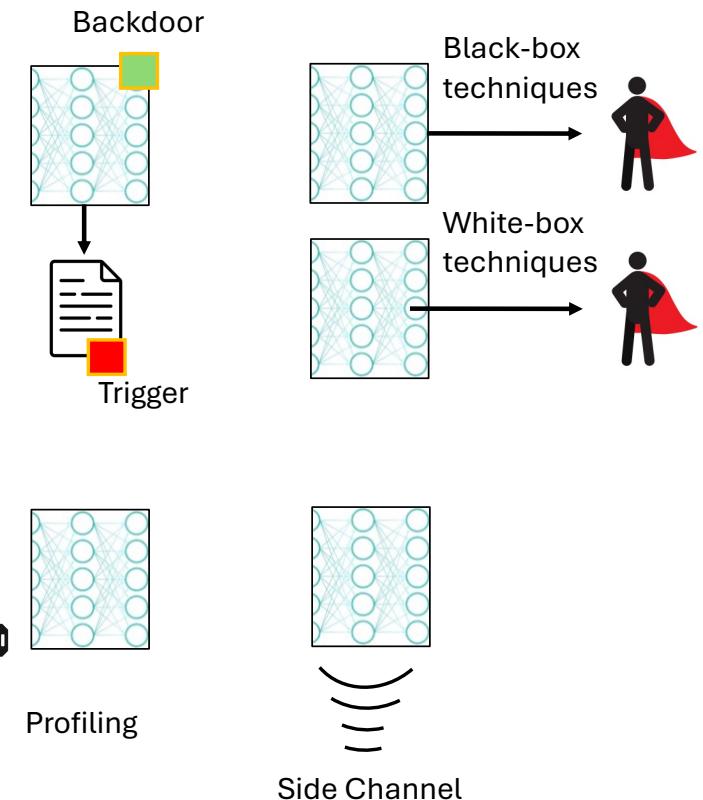
- Some **facts** on the machine learning **industry** in 2021:
 - it was valued \$15.5B
 - the Netflix recommendation engine saves Netflix \$1 billion per year
- **Watermarking AI models is an urgent need:**
 - ...try to do a Google Search
 - many recent works still on Arxiv!
 - funded project



WHAM! - Watermarking Hazards and novel perspectives in Adversarial Machine learning (PRIN - B53D23013340006)

AI Models (2)

- Various techniques:
 - backdoors
 - profiling
 - side channels
 - white/black-box
 - ...
- Example of a **white-box** technique:
 - it uses LSB (again: it is ubiquitous...)
 - it directly embeds data in weights
 - a bit naïve, just an example!
 - a more sophisticated approach is (*)



(*) Y. Nagai, Y Uchida, S. Sakazawa, “Digital Watermarking for Deep Neural Networks”, International Journal of Multimedia Information Retrieval, Vol. 7, pp. 3–16, Feb. 2018

AI Models (2)

- Key idea:

```
[ [ 3.52380052e-02 5.62173128e-01 -1.26857802e-01 -1.08864554e-03  
-7.86279887e-02 9.90515500e-02 1.12150028e-01 2.58520335e-01  
-2.26100925e-02 -1.53759057e-02 -1.05209224e-01 2.85302326e-02  
2.04063114e-02 -1.34602830e-01 2.16201901e-01 -1.56125368e-03  
-2.44291611e-02 4.46301669e-01 -2.84363050e-02 2.77920157e-01  
-3.05880439e-02 1.15929388e-01 -1.89442746e-02 1.00352295e-01  
1.45177217e-02 1.02079324e-01 4.87670563e-02 2.03172565e-01  
-1.09964848e-01 2.30804667e-01 1.12779163e-01 -3.51786196e-01  
2.91943979e-02 1.37668490e-01 -9.48298257e-03 4.48994637e-02  
1.64225727e-01 1.56570017e-01 4.97223698e-02 -1.36244521e-01  
3.86076123e-01 5.40191606e-02 -1.00916103e-02 -1.31050227e-02  
1.42773181e-01 3.65269999e-03 -2.86158137e-02 4.87360284e-02  
-2.37015169e-02 1.14381999e-01 -2.41468489e-01 -6.32218793e-02  
1.30570263e-01 1.62954698e-03 8.21532086e-02 7.68718868e-03  
-5.01250289e-03 8.46664142e-03 -1.45577103e-01 -1.19108362e-02  
1.95129693e-01 3.74974944e-02 2.05875598e-02 1.02373719e-01]  
[ 1.75550357e-02 -1.59322411e-01 -1.64139122e-02 1.46152833e-02  
-1.38553336e-01 2.08842367e-01 1.10801846e-01 1.58453420e-01  
-3.14325979e-03 -8.92286561e-03 2.02904910e-01 1.08194217e-01  
7.81989023e-02 1.67696640e-01 -8.54833331e-03 9.42505430e-03  
-1.28925741e-01 -2.15639502e-01 2.11874619e-02 2.37909436e-01  
-8.93226918e-03 -2.53099948e-01 1.93967856e-02 4.51776162e-02  
-9.96187236e-03 1.14028394e-01 2.00887378e-02 -1.52492542e-02  
5.21601550e-02 -1.98008064e-02 -9.82316434e-02 3.08643609e-01  
-1.11380173e-02 -8.06702003e-02 -1.22004651e-01 1.13300495e-01
```

Original Weights
(Trained Model)

Watermark

The diagram illustrates the process of watermarking. At the top, the text "Watermark" is centered above a downward-pointing arrow. Below the arrow is a rectangular box containing a grid of numerical values. A portion of this grid is highlighted with a pink background, representing the watermark being embedded into the original weights.

```
[ [ 8.80950131e-03 5.62173128e-01 -1.26857802e-01 -1.08864554e-03  
-1.96569972e-02 2.47628875e-01 2.80375071e-02 1.03408134e+00  
-2.26100925e-02 -6.15036227e-01 -2.63023060e-02 2.85302326e-02  
8.16252455e-02 -1.34602830e-01 2.16201901e-01 -6.24501472e-03  
-2.44291611e-02 4.46301669e-01 -1.13745220e-01 1.11168063e+00  
-3.05880439e-02 1.15929388e-01 -1.89442746e-02 2.50880737e-02  
1.45177217e-02 1.02079324e-01 1.21917641e-02 8.12690258e-01  
-2.74912119e-02 2.30804667e-01 1.12779163e-01 -1.40714478e+00  
2.91943979e-02 5.50673962e-01 -9.48298257e-03 4.48994637e-02  
1.64225727e-01 1.56570017e-01 4.97223698e-02 -5.44978082e-01  
3.86076123e-01 5.40191606e-02 -1.00916103e-02 -1.31050227e-02  
1.42773181e-01 3.65269999e-03 -2.86158137e-02 4.87360284e-02  
-2.37015169e-02 1.14381999e-01 -2.41468489e-01 -6.32218793e-02  
1.30570263e-01 1.62954698e-03 8.21532086e-02 7.68718868e-03  
-5.01250289e-03 8.46664142e-03 -1.45577103e-01 -1.19108362e-02  
1.95129693e-01 3.74974944e-02 2.05875598e-02 1.02373719e-01]  
[ 1.75550357e-02 -1.59322411e-01 -1.64139122e-02 1.46152833e-02  
-1.38553336e-01 2.08842367e-01 1.10801846e-01 1.58453420e-01  
-3.14325979e-03 -8.92286561e-03 2.02904910e-01 1.08194217e-01  
7.81989023e-02 1.67696640e-01 -8.54833331e-03 9.42505430e-03  
-1.28925741e-01 -2.15639502e-01 2.11874619e-02 2.37909436e-01  
-8.93226918e-03 -2.53099948e-01 1.93967856e-02 4.51776162e-02  
-9.96187236e-03 1.14028394e-01 2.00887378e-02 -1.52492542e-02  
5.21601550e-02 -1.98008064e-02 -9.82316434e-02 3.08643609e-01  
-1.11380173e-02 -8.06702003e-02 -1.22004651e-01 1.13300495e-01
```

Watermarked Weights
(Trained Model)

Examples/AI

AI Models (2)

- Key idea:

```

[[ 3.52380052e-02 5.62173128e-01 -1.26857802e-01 -1.08864554e-03
-7.86279887e-02 9.90515500e-02 1.12150028e-01 2.58520335e-01
-2.26100925e-02 -1.53759057e-02 -1.05209224e-01 2.85302326e-02
2.04063114e-02 -1.34602830e-01 2.16201901e-01 -1.56125368e-03
-2.44291611e-02 4.46301669e-01 -2.84363050e-02 2.77920157e-01
-3.05880439e-02 1.15929388e-01 -1.89442746e-02 1.00352295e-01
1.45177217e-02 1.02079324e-01 4.87670563e-02 2.03172565e-01
-1.09964848e-01 2.30804667e-01 1.12779163e-01 -3.51786196e-01
2.91943979e-02 1.37668490e-01 -9.48298257e-03 4.48994637e-02
1.64225727e-01 1.56570017e-01 4.97223698e-02 -1.36244521e-01
3.86076123e-01 5.40191606e-02 -1.00916103e-02 -3.13050227e-02
1.42773181e-01 3.65269999e-03 -2.86158137e-02 4.87360284e-02
-2.37015169e-02 1.14381999e-01 -2.41468489e-01 -6.32218793e-02
1.30570263e-01 1.62954698e-03 8.21532086e-02 7.68718868e-03
-5.01250289e-03 8.46664142e-03 -1.45577103e-01 -1.19108362e-02
1.95129693e-01 3.74974944e-02 2.05875598e-02 1.02373719e-01]
[ 1.75550337e-02 -1.59322411e-01 -1.64139122e-02 1.46125833e-02
-1.38553336e-01 2.08842367e-01 1.10801846e-01 1.58453420e-01
-3.14325979e-03 -8.92286561e-03 2.02904910e-01 1.08194217e-01
7.81989023e-02 1.67696640e-01 -8.54833331e-03 9.42505430e-03
-1.28925741e-01 -2.15639502e-01 2.11874619e-02 2.37909436e-01
-8.93226918e-03 -2.53099948e-01 1.93967856e-02 4.51776162e-02
-9.96187236e-03 1.14028394e-01 2.00887378e-02 -1.52492542e-02
5.21601559e-02 -1.98080646e-02 -9.82316434e-02 3.08643609e-01
-1.11380173e-02 -8.06702003e-02 -1.22004651e-01 1.13300495e-01

```

Original Weights (Trained Model)

MNIST Dataset

Accuracy: 98.58%

Examples/AI

AI Models (2)

- Key idea:

[8.80950131e-03	5.62173128e-01	-1.26857802e-01	-1.08864554e-03
-1.96569972e-02	2.47628875e-02	2.80375071e-02	1.034018134e+00	
-2.26100925e-02	-6.15036227e-02	-2.63023060e-02	2.85302326e-02	
8.16252455e-02	-1.34602830e-01	2.16201901e-01	-6.24501472e-03	
-2.44291611e-02	4.46301669e-01	-1.13745220e-01	1.11168063e+00	
-3.05880439e-02	1.15929388e-01	-1.89442746e-02	2.50880737e-02	
1.45177217e-02	1.02079324e-01	1.21917641e-02	8.12690258e-01	
-2.74912119e-02	2.30804667e-01	1.12779163e-01	-4.07144748e+00	
2.91943979e-02	5.50673962e-01	-9.48298257e-03	4.48994637e-02	
1.64225727e-01	1.56570017e-01	4.97223698e-02	-5.44978082e-01	
3.86076123e-01	5.40191606e-02	-1.00916103e-02	-1.31050227e-02	
1.42773181e-01	3.65269999e-03	-2.86158137e-02	4.87360284e-02	
-2.37015169e-02	1.14381999e-01	-2.41468489e-01	-6.32218793e-02	
1.30570263e-01	1.62954698e-03	8.21523086e-02	7.68718868e-03	
-5.01250289e-03	8.46664142e-03	-1.45577103e-01	-1.19108362e-02	
1.95129693e-01	3.74974944e-02	2.05875598e-02	1.02373719e-01	
[1.75550357e-02	-1.59322411e-01	-1.64139122e-02	1.46152833e-02
-1.38553336e-01	2.08842367e-01	1.10801846e-01	1.58453420e-01	
-3.14325979e-03	8.92286561e-03	2.02964910e-01	1.08194271e-01	
7.81989023e-02	1.67696640e-01	-8.54833331e-03	9.42505430e-03	
-1.28925741e-01	-2.15639502e-01	2.11874619e-02	2.37909436e-01	
-8.93226918e-03	-2.53309948e-01	1.93967856e-02	4.51776162e-02	
-9.96187236e-03	1.14028394e-01	2.00887378e-02	-1.52492542e-02	
5.21601550e-02	-1.98080604e-02	-9.82316434e-02	3.08643609e-01	
-1.11380173e-02	-8.06702003e-02	-1.22004651e-01	1.13300495e-01	

Watermarked Weights (Trained Model)

MNIST Dataset

Accuracy: 94.68%

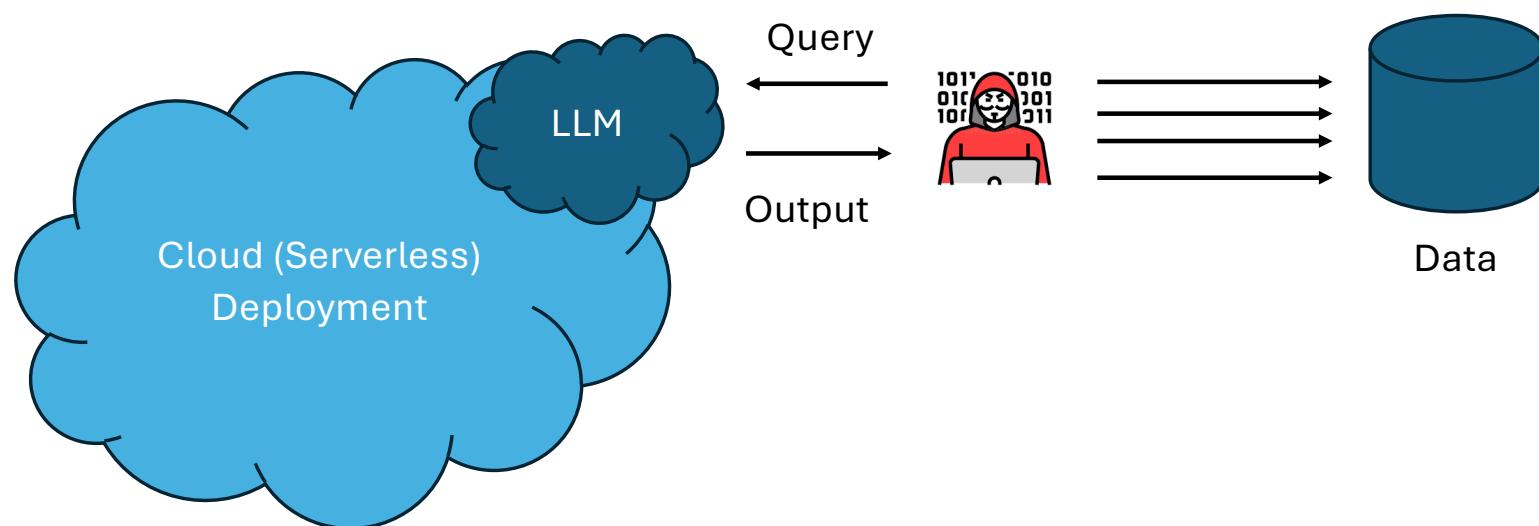
To reduce the impact of the watermark a possible approach is to use a “dropout” mechanism (before the FC2 layer) to induce an ensemble behavior. The accuracy increases again to ~98%!

AI Outputs (1)

- Large Language Models (LLMs) spawned **new services**:
 - chatbots
 - financial analysis
 - content generation
 - code generation
 - ...
- Commercial LLMs offers some form of APIs:
 - to **avoid training**, which is expensive and time-consuming
 - yet, an attacker can perform **multiple queries** to imitate the model
- **We will threat LLM as a synecdoche of as-a-Service frameworks**

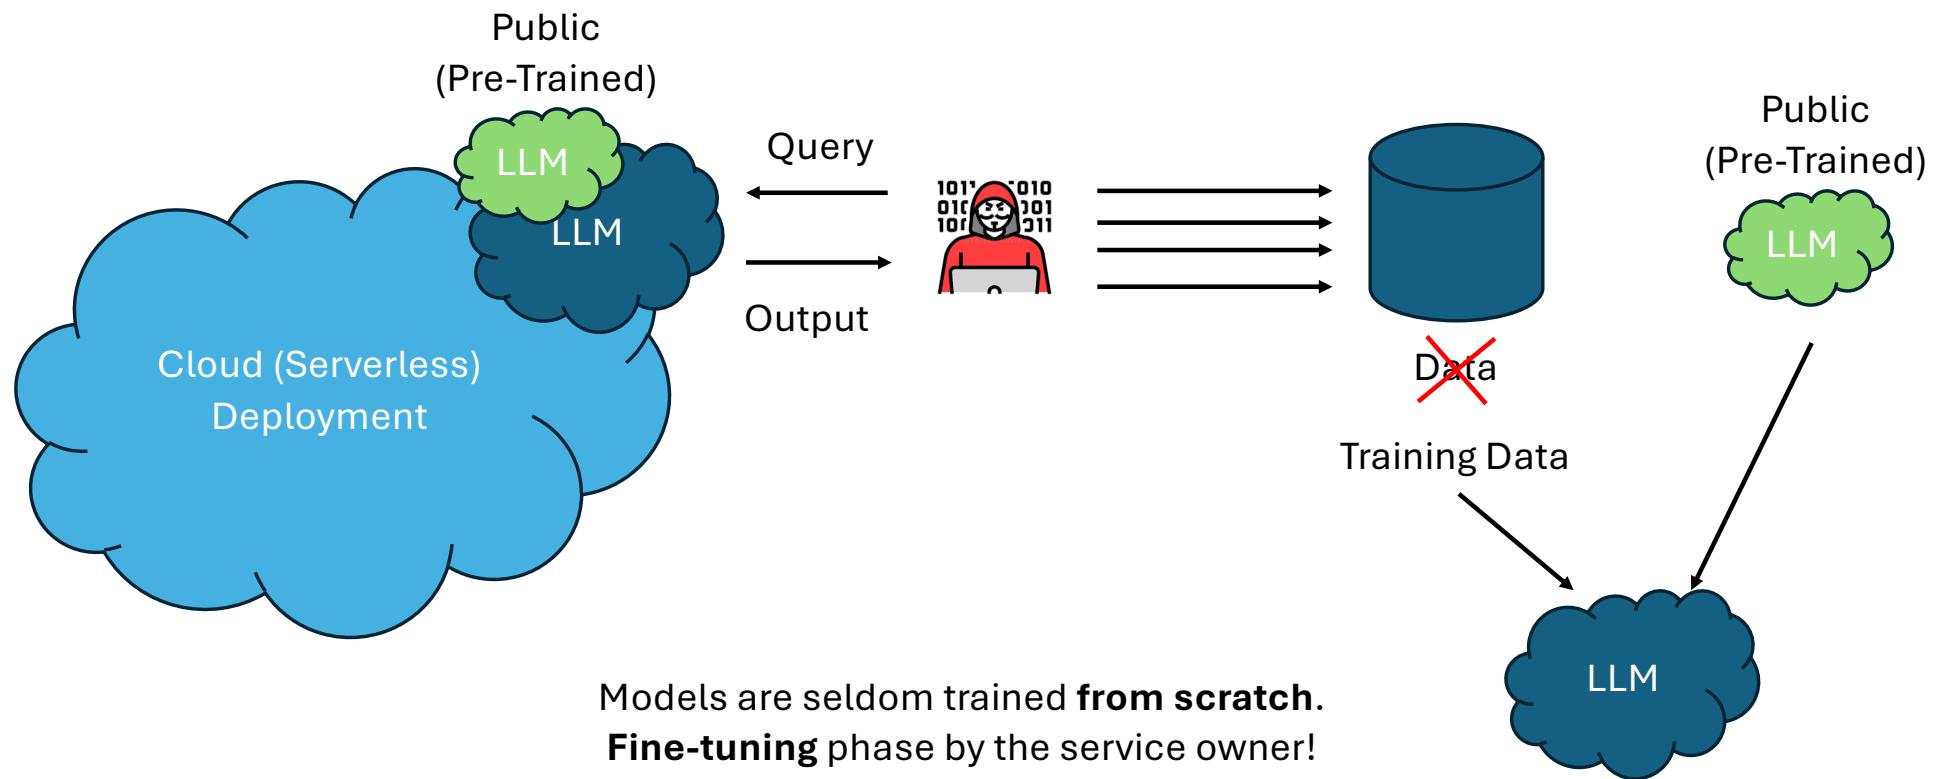
AI Outputs (2)

- General attack model:



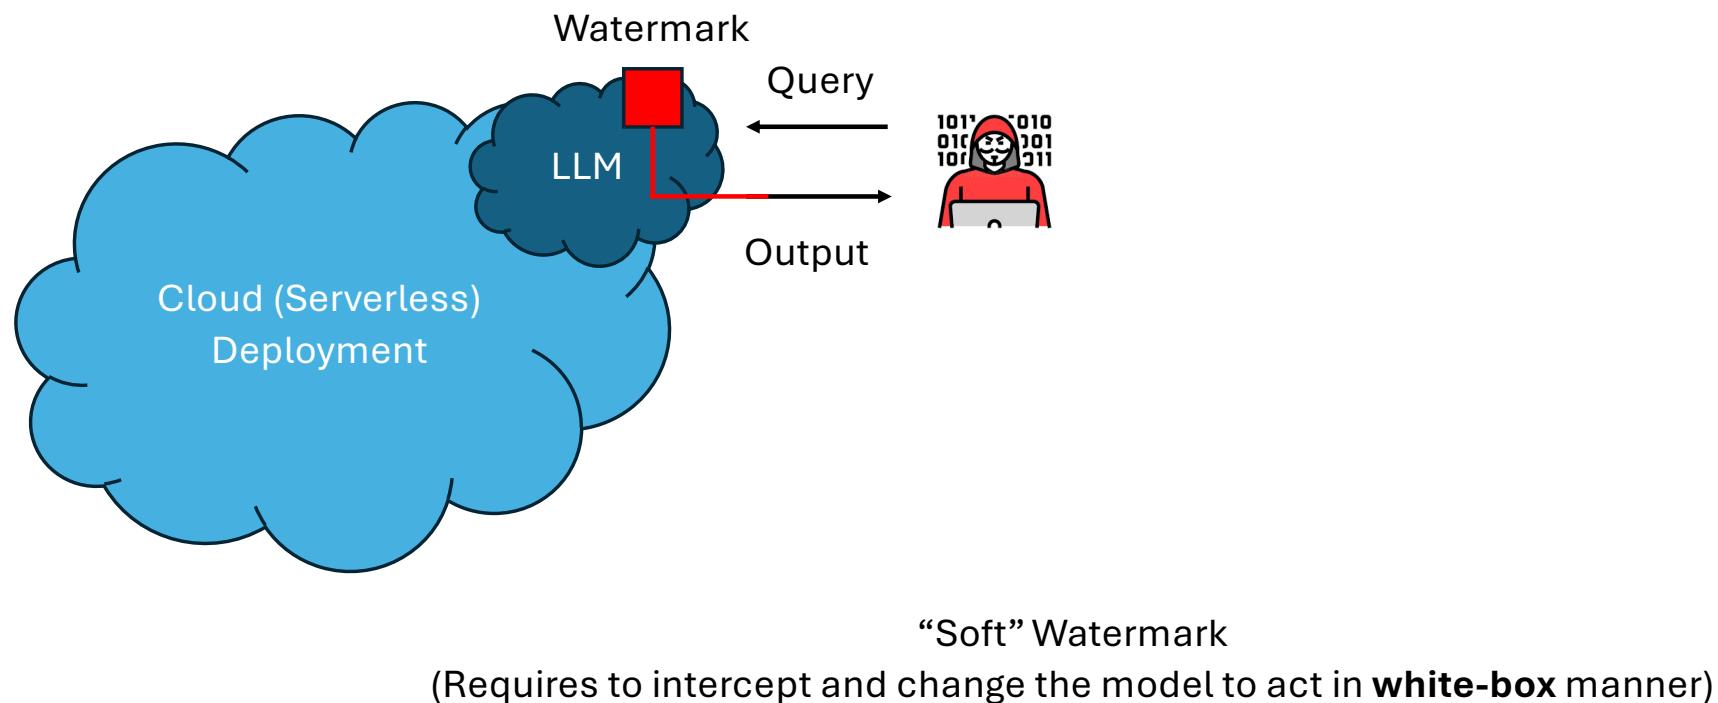
AI Outputs (2)

- General attack model:



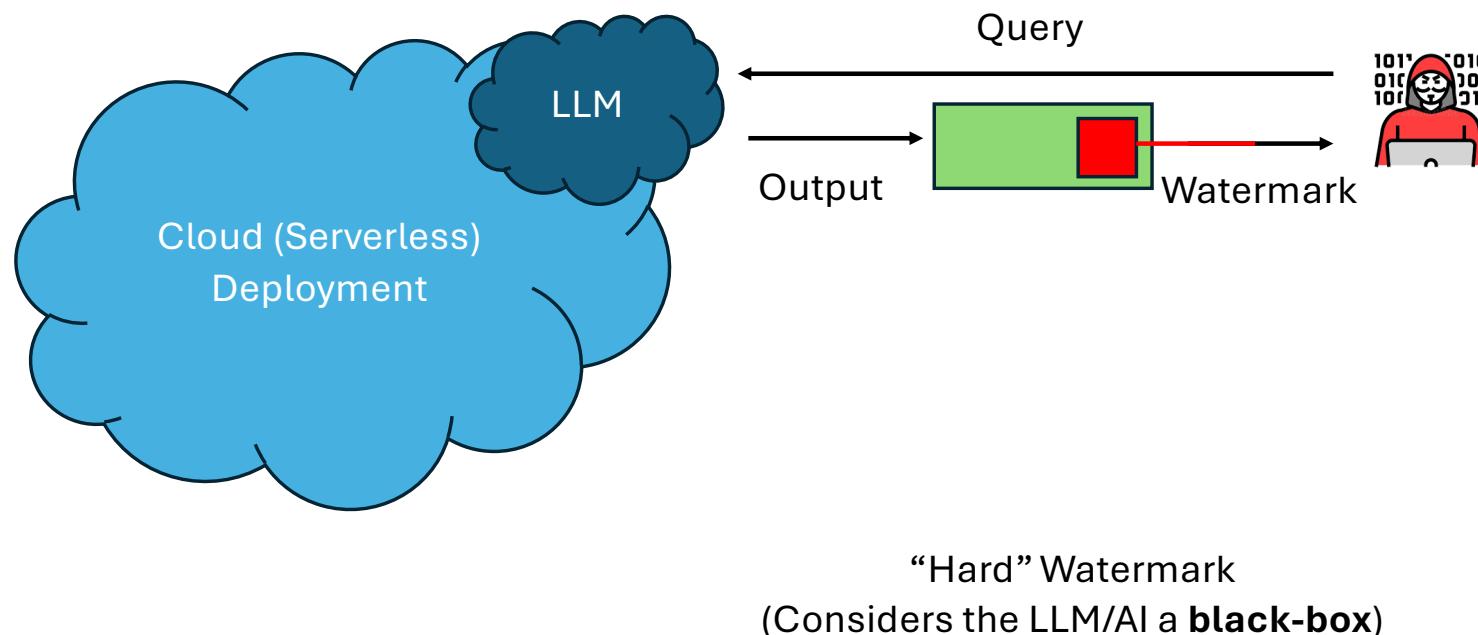
AI Outputs (2)

- General attack model:



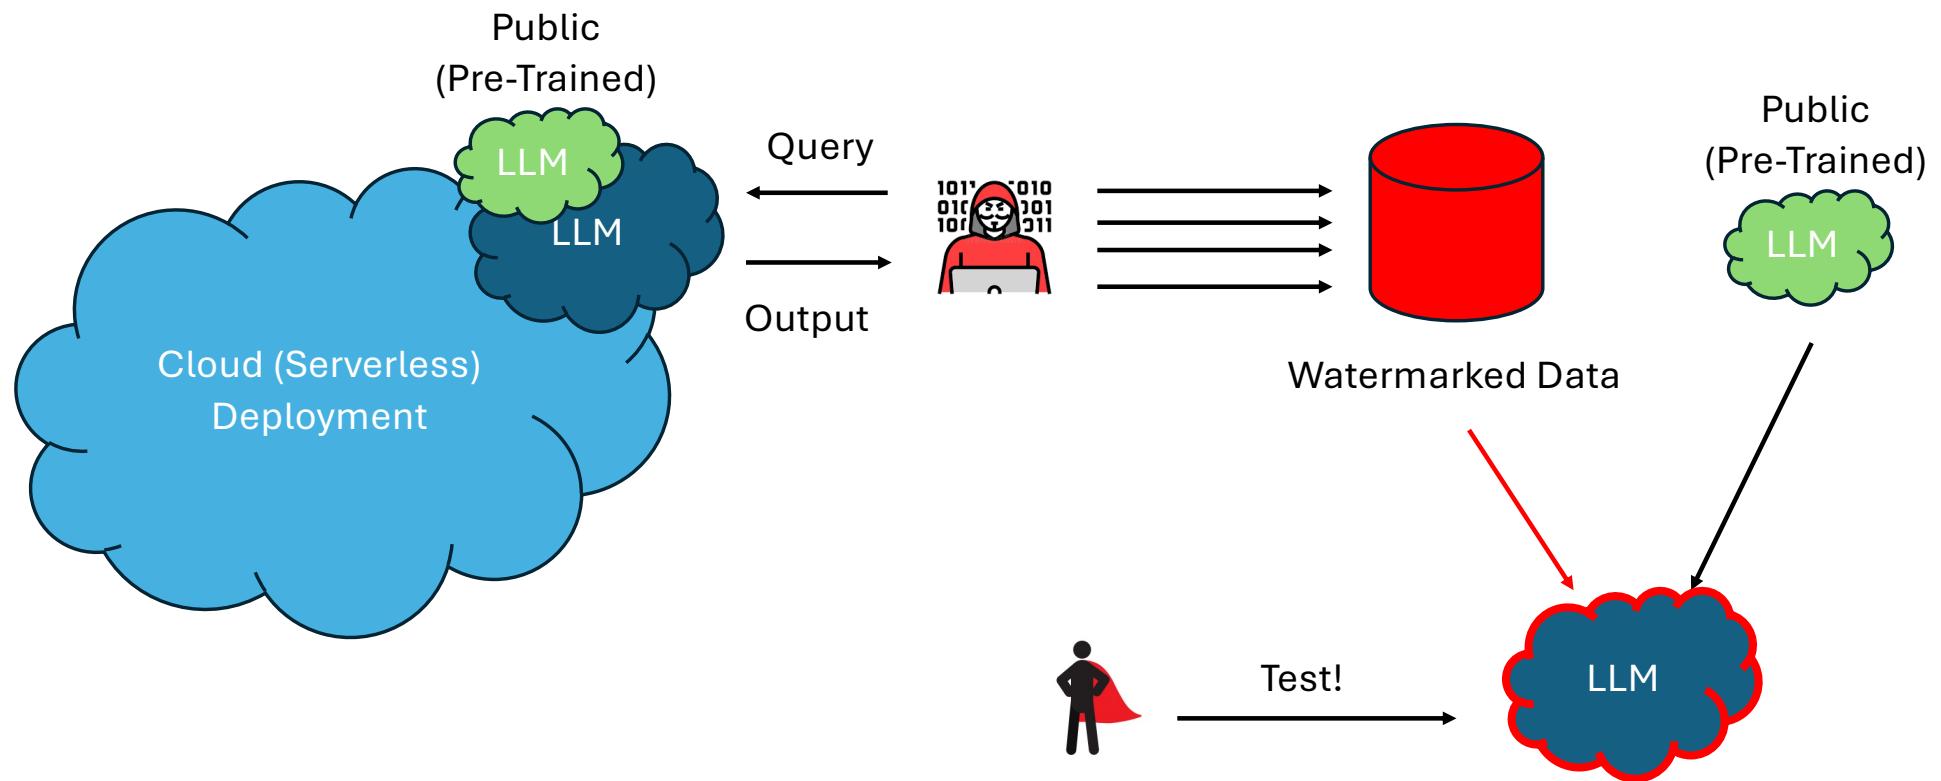
AI Outputs (2)

- General attack model:



AI Outputs (2)

- General attack model:



AI Outputs (3)

- Example of a **black-box** technique:
 - for marking LLM-based code generation tools (very popular!)
 - from Li, Wang, Wang and Gao(*)

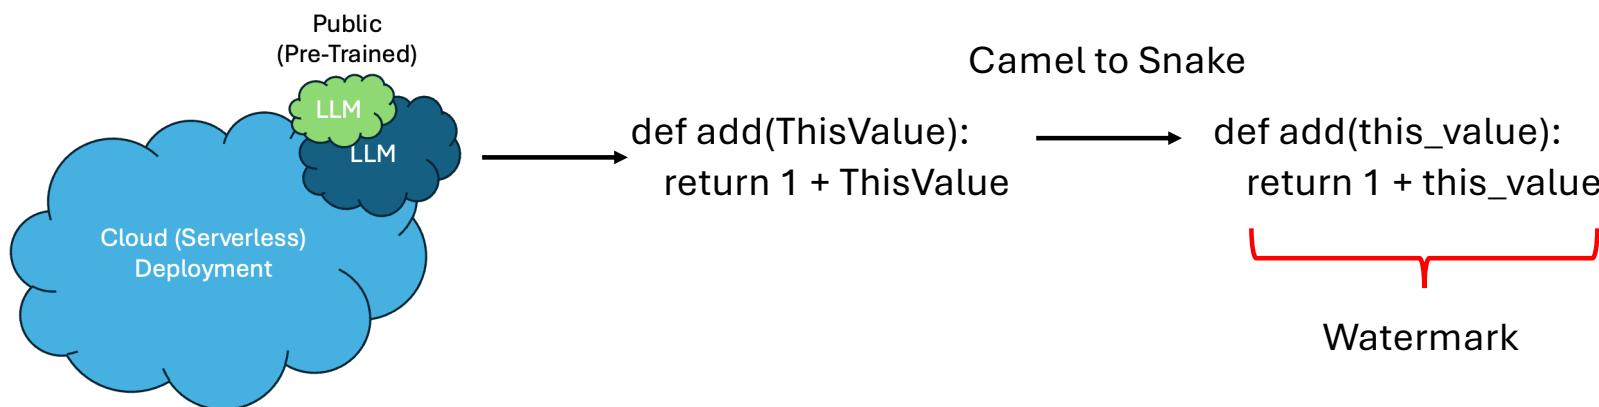
AI Outputs (3)

- Example of a **black-box** technique:
 - for marking LLM-based code generation tools (very popular!)
 - from Li, Wang, Wang and Gao(*)
- Key idea:
 - AI-generated code is changed by replacing tokens with synonyms
 - the process is iterated to match a target distribution (for test)
 - many modern programming languages can be altered without changing their semantics



AI Outputs (3)

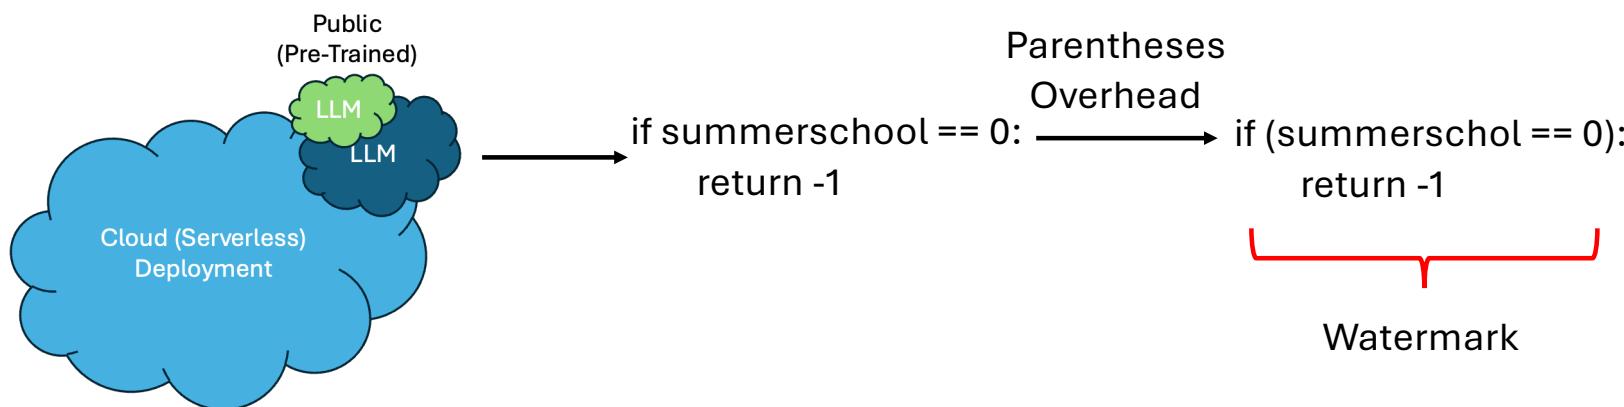
- Example of a **black-box** technique:
 - for marking LLM-based code generation tools (very popular!)
 - from Li, Wang, Wang and Gao(*)
- Key idea:



(*) Z. Li, C. Wang, S. Wang, C. Gao, "Protecting Intellectual Property of Large Language Model-Based Code Generation APIs via Watermarks", Proc. of the ACM SIGSAC Conference on Computer and Communications Security, pp. 2336 – 2350, Nov. 2023

AI Outputs (3)

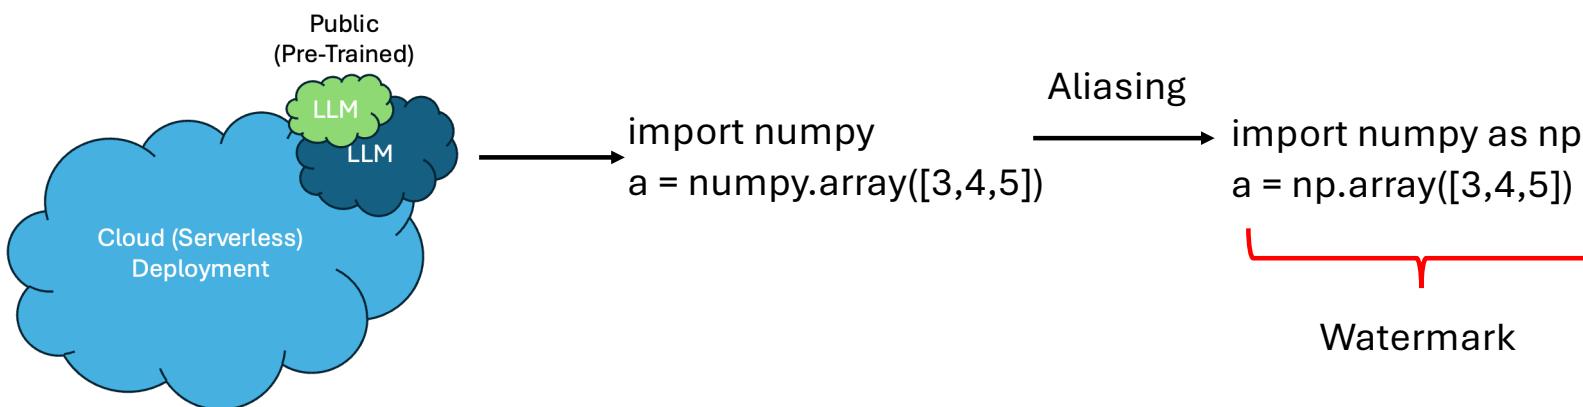
- Example of a **black-box** technique:
 - for marking LLM-based code generation tools (very popular!)
 - from Li, Wang, Wang and Gao(*)
- Key idea:



(*) Z. Li, C. Wang, S. Wang, C. Gao, “Protecting Intellectual Property of Large Language Model-Based Code Generation APIs via Watermarks”, Proc. of the ACM SIGSAC Conference on Computer and Communications Security, pp. 2336 – 2350, Nov. 2023

AI Outputs (3)

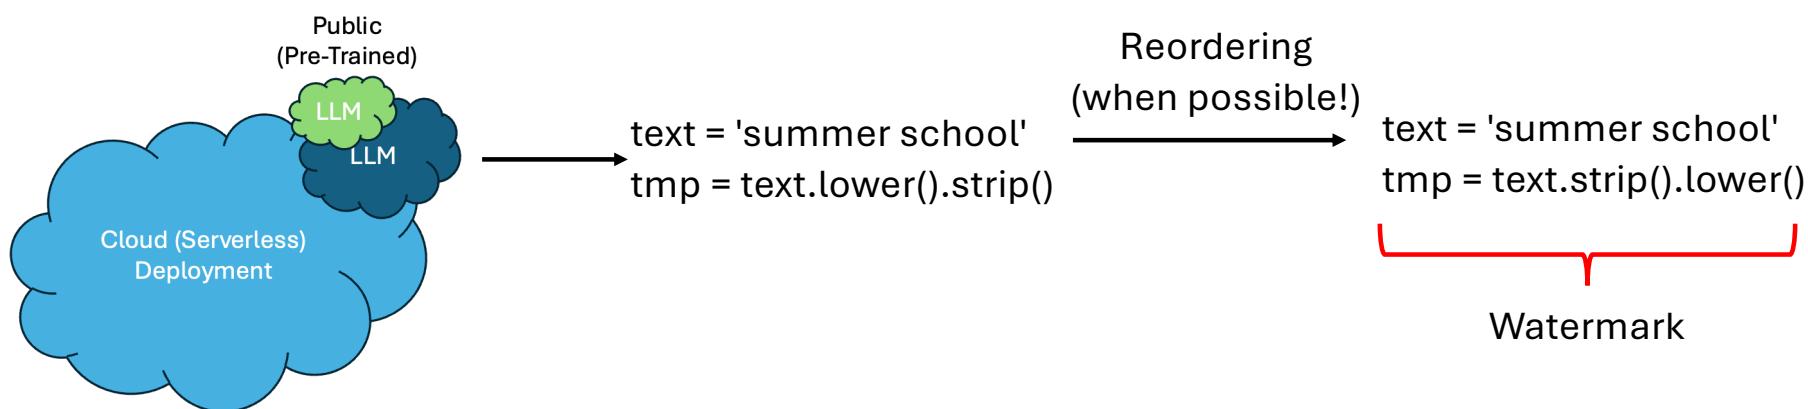
- Example of a **black-box** technique:
 - for marking LLM-based code generation tools (very popular!)
 - from Li, Wang, Wang and Gao(*)
- Key idea:



(*) Z. Li, C. Wang, S. Wang, C. Gao, “Protecting Intellectual Property of Large Language Model-Based Code Generation APIs via Watermarks”, Proc. of the ACM SIGSAC Conference on Computer and Communications Security, pp. 2336 – 2350, Nov. 2023

AI Outputs (3)

- Example of a **black-box** technique:
 - for marking LLM-based code generation tools (very popular!)
 - from Li, Wang, Wang and Gao(*)
- Key idea:



(Some) Attacks

Attacks (1)

- Attacks heavily depend on:
 - the **type** of watermark
 - the protected media/content, i.e., the **carrier**
- Manipulations could be **intentional** or **unintentional**
- Several **attack models** are possible:
 - different degree of sophistication
 - exploiting diverse attack surfaces
 - tweaked for specific contexts, e.g., AI vs cloud
- Attacks can be “chained” to implement advanced **offensive templates**



Running out of time: attack models for many AI scenarios have been nicely described in: F. Boenisch, “A Systematic Review on Model Watermarking for Neural Networks”, *Frontiers in Big Data*, Vol. 4, pp. 4:1 – 4:16, Nov. 2021

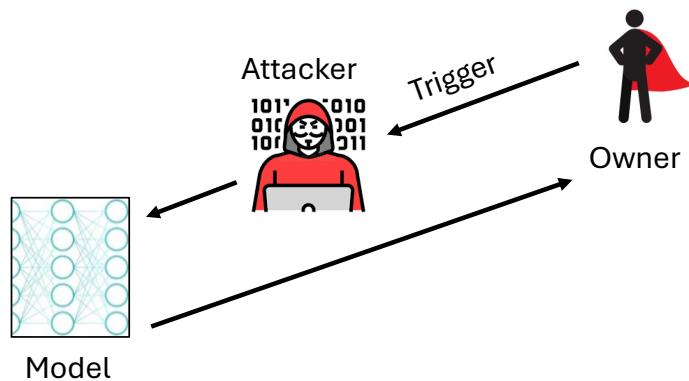
Attacks (2)

- The most important attack models are:
 - **detection**: attacker can discover the presence of a watermark

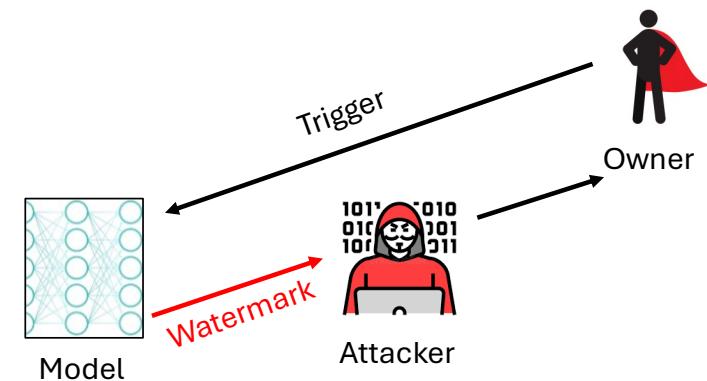


Attacks (2)

- The most important attack models are:
 - detection: attacker can discover the presence of a watermark
 - **suppression**: attacker can impede to access the watermark



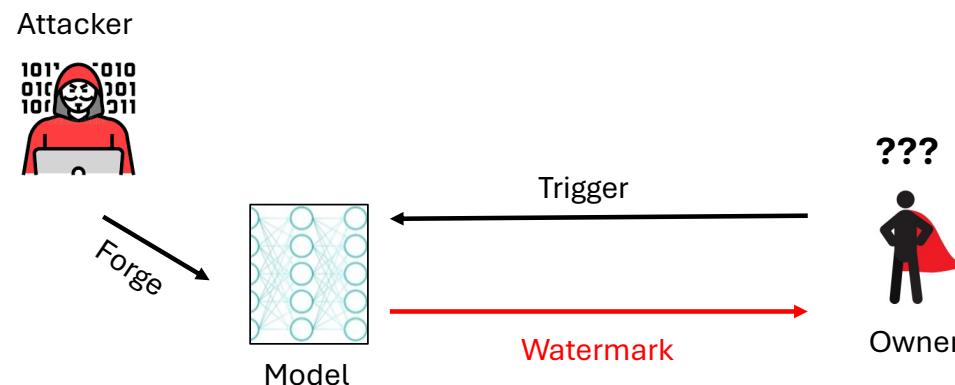
Trigger Suppression



Watermark Suppression

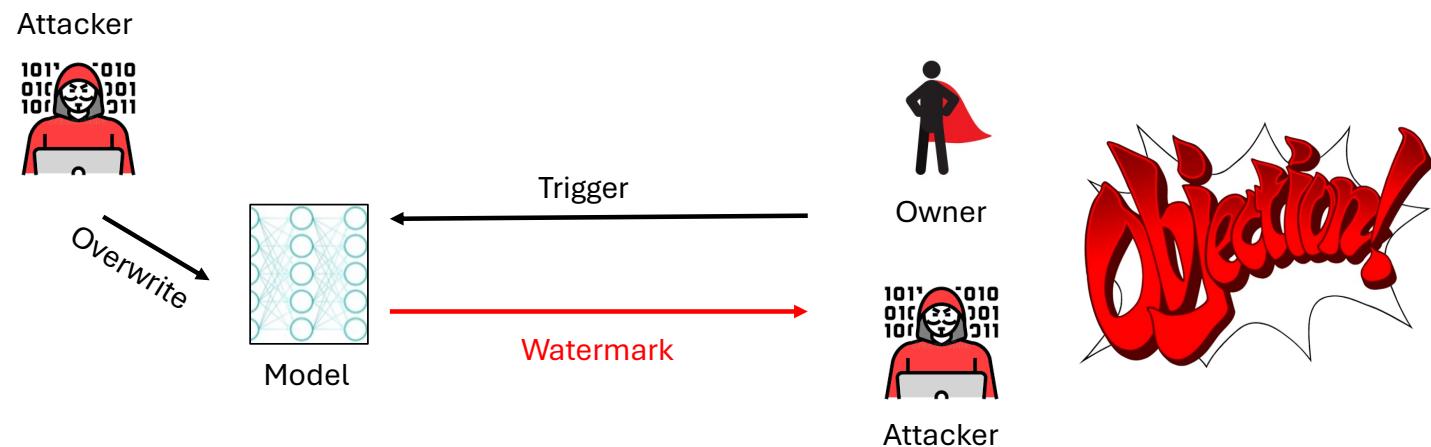
Attacks (2)

- The most important attack models are:
 - detection: attacker can discover the presence of a watermark
 - suppression: attacker can impede to access the watermark
 - **forging**: attacker can render the watermark ambiguous



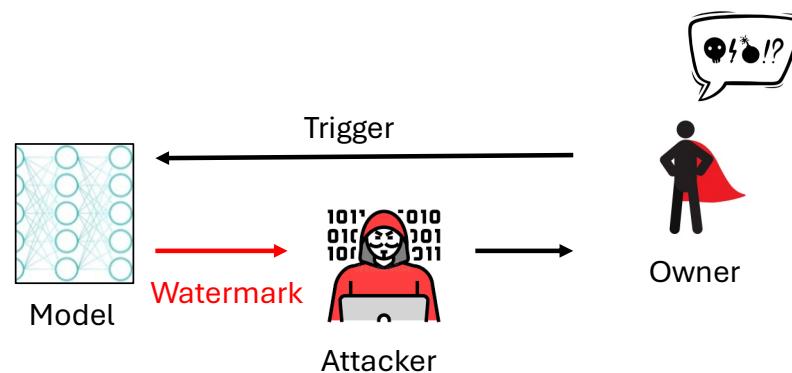
Attacks (2)

- The most important attack models are:
 - detection: attacker can discover the presence of a watermark
 - suppression: attacker can impede to access the watermark
 - forging: attacker can render the watermark ambiguous
 - **overwrite**: attacker can substitute the watermark to pretend ownership



Attacks (2)

- The most important attack models are:
 - detection: attacker can discover the presence of a watermark
 - suppression: attacker can impede to access the watermark
 - forging: attacker can render the watermark ambiguous
 - overwrite: attacker can substitute the watermark to pretend ownership
 - **removal**: attacker can prevent the use of watermark

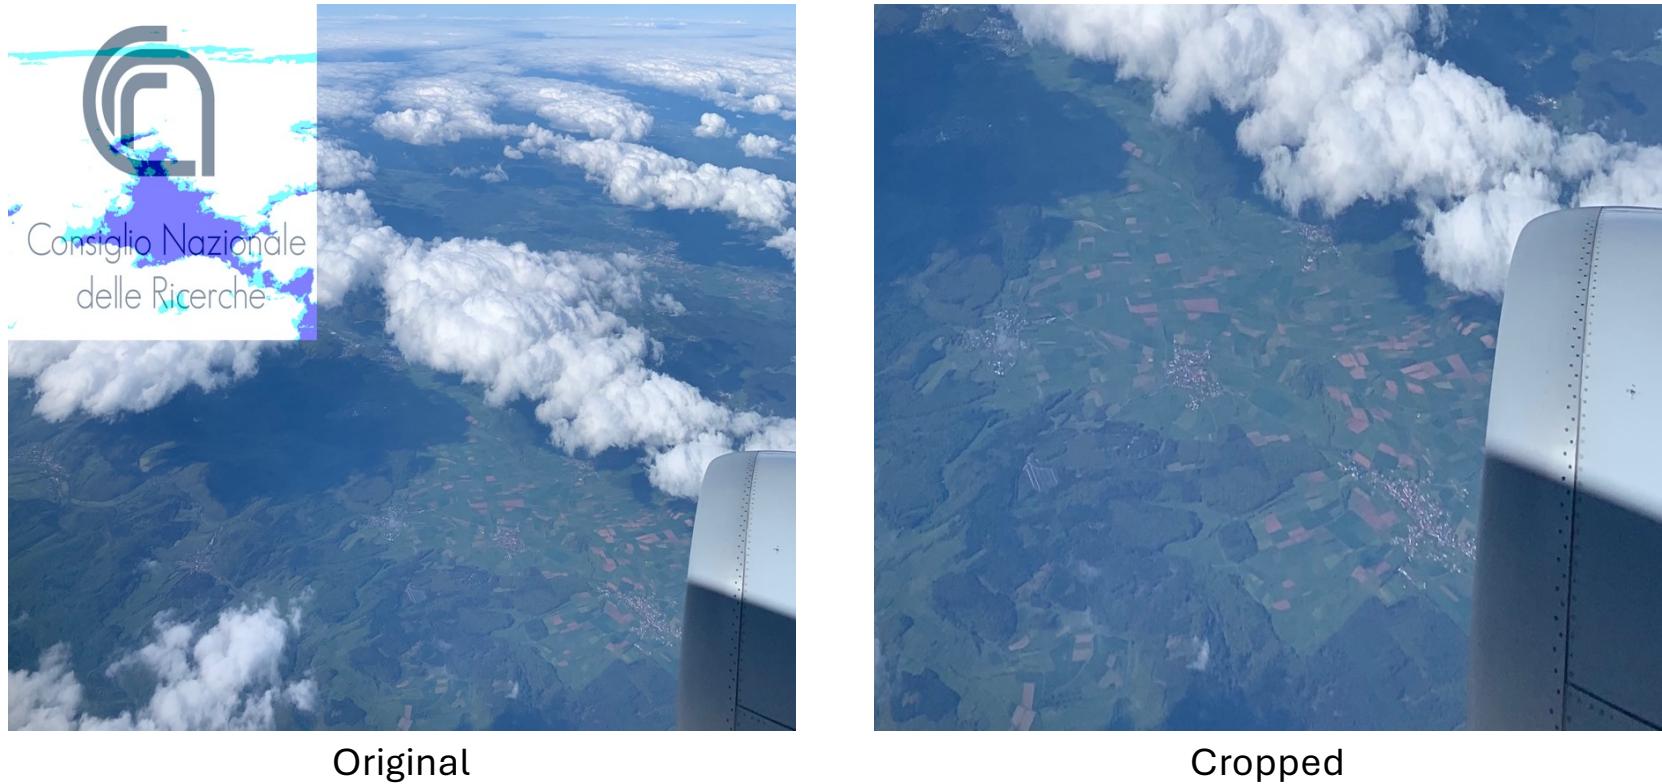


Attacks (3)

- Removal attacks are the most popular, especially for:
 - multimedia
 - digital objects
 - 3D mesh
 - network traffic
- The main classes of **attacks** are:
 - additive noise
 - (pseudo) random alterations
 - remove frequencies via filtering
 - cropping, rotation, and scaling
 - smoothing and simplification
 - compression or transcoding
 - statistical averaging
 - multiple watermarking
 - spoofing
 - ...

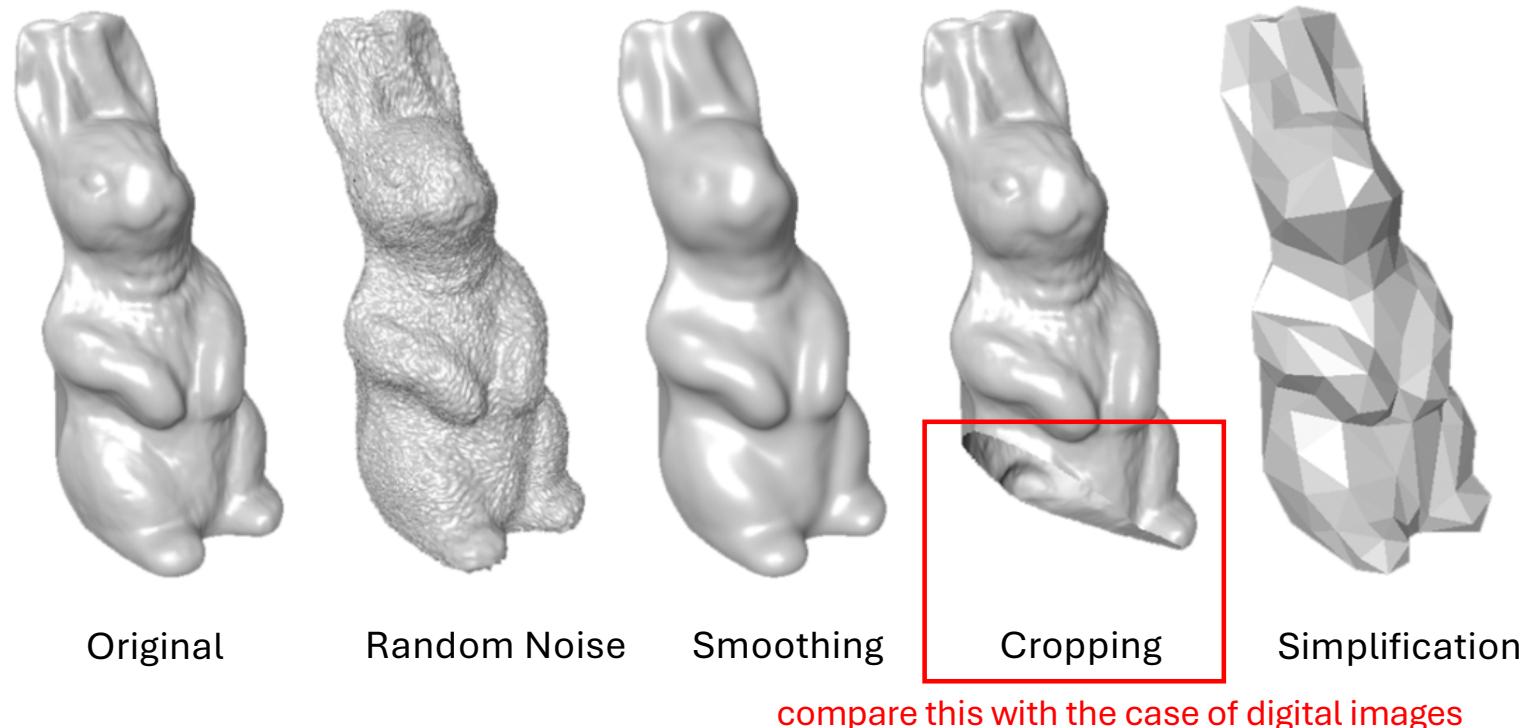
Attacks (4)

- Example 1 (cropping attack against a digital image):



Attacks (4)

- Example 2 (various attacks against a 3D mesh):



compare this with the case of digital images

Security Implications

Security Implications (1)

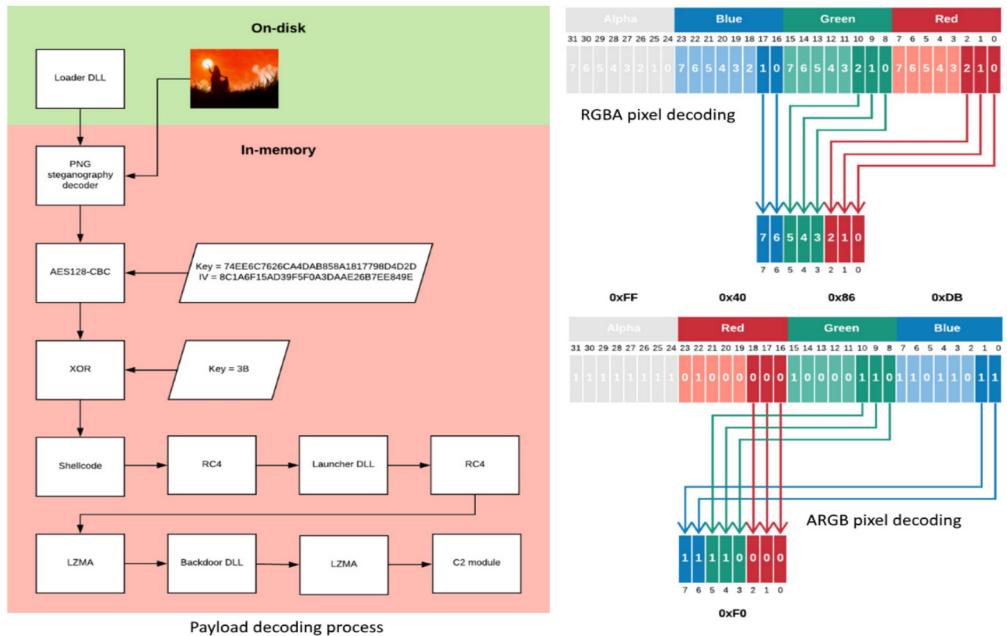
- Watermarking is **tightly-coupled** with:
 - information hiding
 - steganography
- If you can **hide data** for IP, fingerprinting, and integrity **you can**:
 - track (unaware) individuals
 - hide malicious contents
 - smuggle secrets
 - industrial espionage
 - ...
- In essence:
 - **opportunity makes the thief!**

Security Implications (2)

- Despite the **effort** of many **security experts** and **researchers**:
 - **countermeasures** are progressively showing **limitations**
 - only a **fraction** of threats is **detected**
 - malware increasingly operates **undisturbed** for **longer timeframes**
- How can malware developers **avoid detection** for long periods?
- Giving a final **answer** is very **hard**, but:
 - steganography and information hiding
 - covert channels
- Attackers usually borrows good ideas from other fields:
 - **watermarking techniques** should be considered double-edged

Security Implications (3)

- Key idea under **stegomalware** (*steganographic malware*):



OceanLotus(*) uses
steganography

Source for OceanLotus: <https://www.bleepingcomputer.com/news/security/oceanlotus-apt-uses-steganography-to-load-backdoors/>

(*) L. Caviglione, W. Mazurczyk, "Never Mind the Malware, Here's the Stegomalware", IEEE Security & Privacy, vol. 20, no. 5, pp. 101-106, Sept.-Oct. 2022

Security Implications (4)

- Example of OceanLotus (*encrypted payload hidden in a .png*):



Original Image

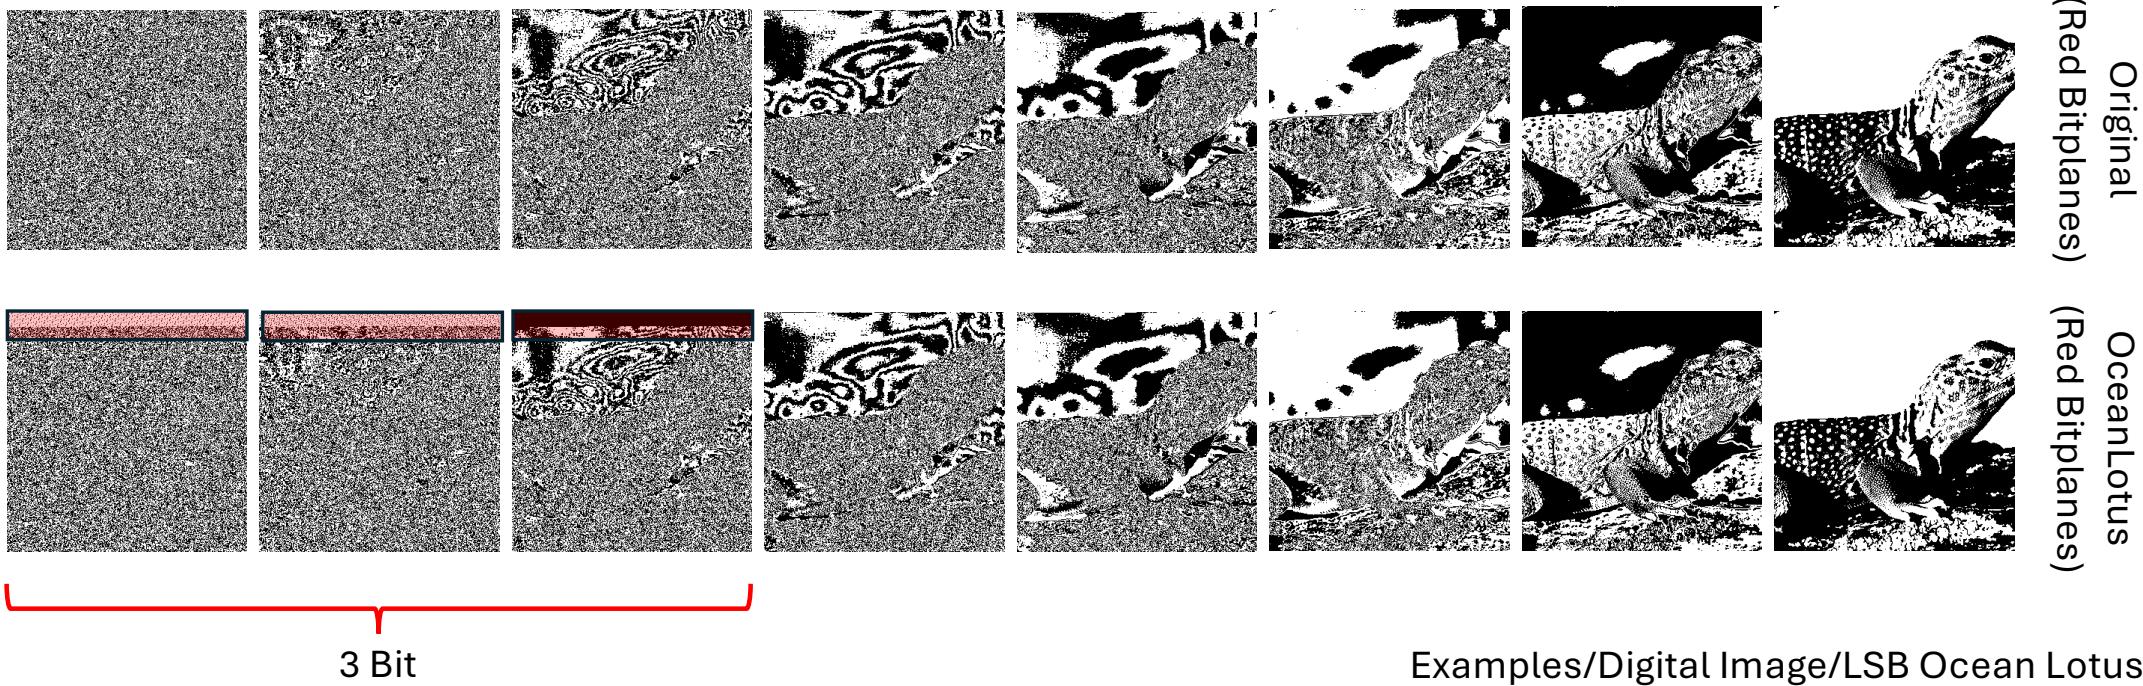


OceanLotus Image

Examples/Digital Image/LSB Ocean Lotus

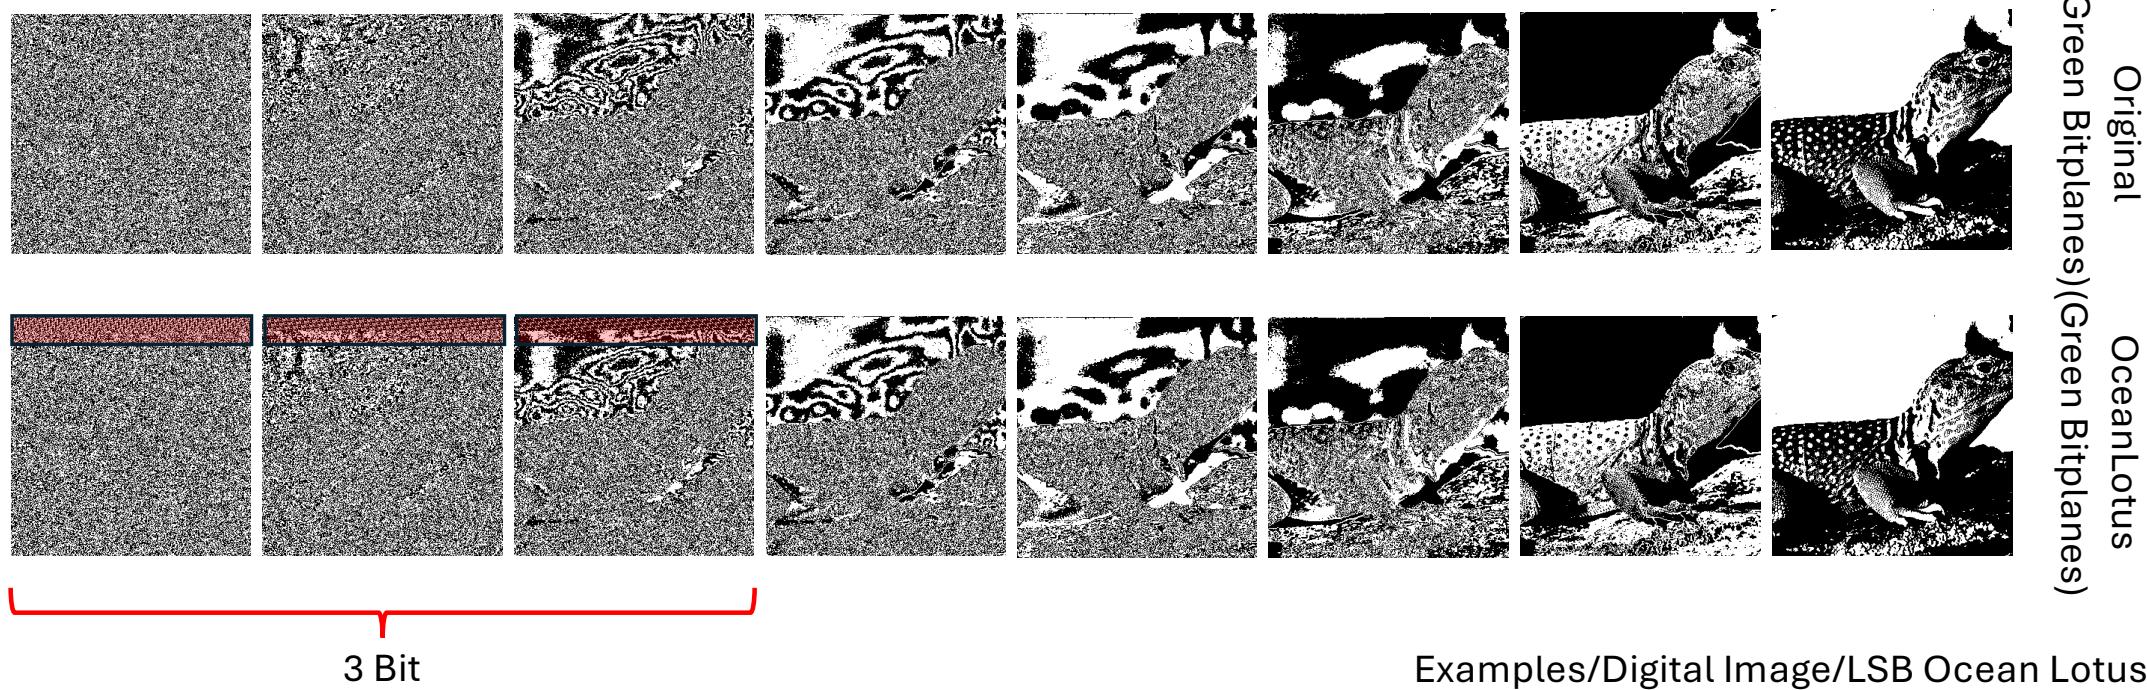
Security Implications (4)

- Example of OceanLotus (*encrypted payload hidden in a .png*):



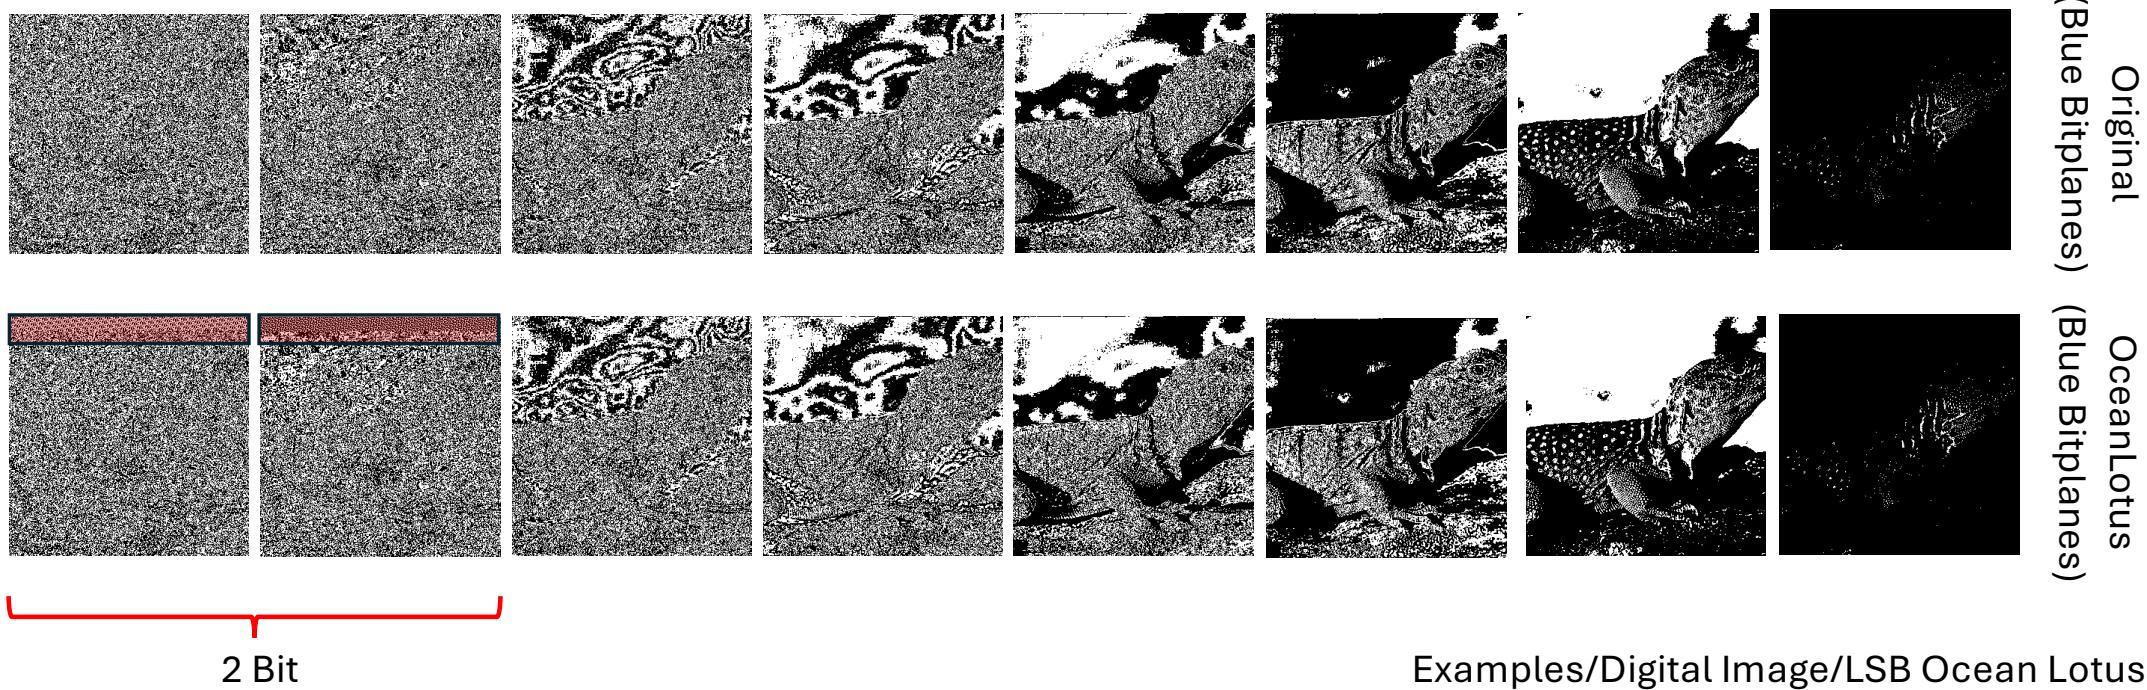
Security Implications (4)

- Example of OceanLotus (*encrypted payload hidden in a .png*):



Security Implications (4)

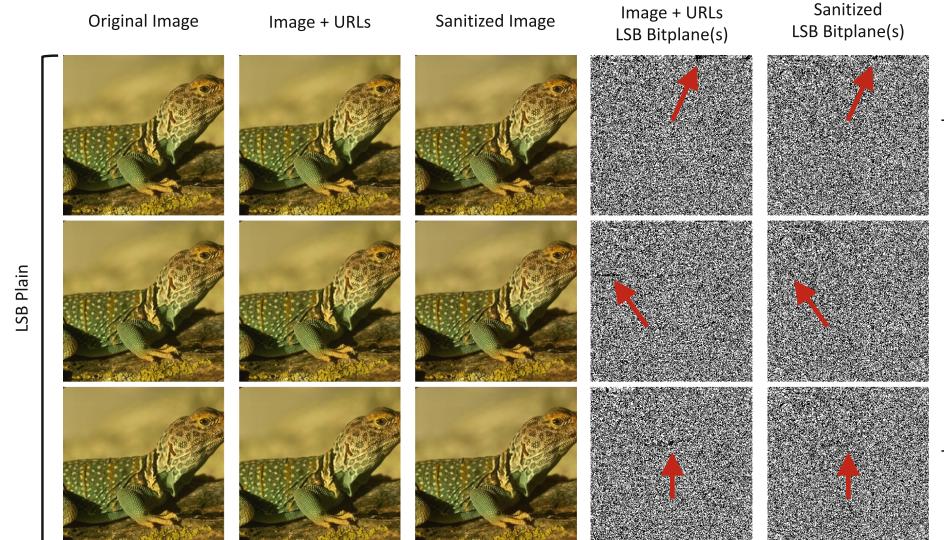
- Example of OceanLotus (*encrypted payload hidden in a .png*):



Open Research Challenges

Open Research Challenges (1)

- **Prevent** that a method could be used to **conceal malicious data**:
 - try to hide the minimum amount of information needed
- **Discriminate** between “good” and “bad” content:
 - already “pretty good” at detecting and sanitizing stegomalware with AI(*)



Note: “sanitizing” an image containing some steganographic data is very close to attacking a watermark!

Open Research Challenges (1)

- **Prevent** that a method could be used to **conceal malicious data**:
 - try to hide the minimum amount of information needed
- **Discriminate** between “good” and “bad” content:
 - already “pretty good” at detecting and sanitizing stegomalware with AI(*)
- This is (again!) a **double-edge** sword:
 - good **sanitizers** could be also effective in **removing watermarks**
 - valuable for **attackers!**
- Create “**selective**” sanitizers:
 - malware is removed, watermark is preserved

Open Research Challenges (2)

- Watermarking techniques can be applied to **AI** to create a **new-wave** of stegomalware

Open Research Challenges (2)

- Watermarking techniques can be applied to **AI** to create a **new-wave** of stegomalware
- A malicious **payload** can be **cloaked** via(*):
 - **classical LSB** approaches, i.e., the last bit(s) of some model parameters are altered to encode arbitrary data



(*) T. Liu, et al., "Stegonet: Turn Deep Neural Network into a Stegomalware", Proc. of the Annual Computer Security Applications Conference, New York, NY, USA, 2020, p. 928–938

Open Research Challenges (2)

- Watermarking techniques can be applied to **AI** to create a **new-wave** of stegomalware
- A malicious **payload** can be **cloaked** via(*):
 - **exploiting resilience** of deep neural networks to introduce internal errors by overwriting the model with secret data and do not re-train the resulting “broken neurons”



(*) T. Liu, et al., "Stegonet: Turn Deep Neural Network into a Stegomalware", Proc. of the Annual Computer Security Applications Conference, New York, NY, USA, 2020, p. 928–938

Open Research Challenges (2)

- Watermarking techniques can be applied to **AI** to create a **new-wave** of stegomalware
- A malicious **payload** can be **cloaked** via(*):
 - **mapping** techniques, i.e., arbitrary information is placed by altering parameters with similar or closest values



(*) T. Liu, et al., "Stegonet: Turn Deep Neural Network into a Stegomalware", Proc. of the Annual Computer Security Applications Conference, New York, NY, USA, 2020, p. 928–938

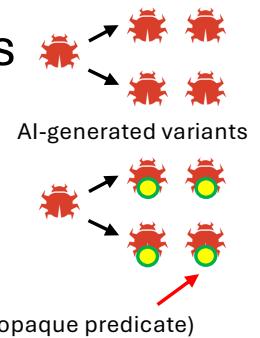
Open Research Challenges (3)

- Watermarking techniques are in general **very specialized**, but:
 - based on **recurring** operations (e.g., moving, reordering, and delaying)
 - sometimes with a good level of **abstraction** (e.g., LSB)
- Finding some **embedding patterns**(*) is desirable to:
 - pursue optimization
 - anticipate attacks or misuses observed in similar cloaking algorithms
 - create general templates
 - fix the terminology and unify the literature
- A major optimization concerns the **green-AI**:
 - watermarking large-scale datasets could be computing/energy consuming

(*) S. Wendzel, L. Caviglione, W. Mazurczyk, A. Mileva, J. Dittmann, C. Krätzer, K. Lamshöft, C. Vielhauer, L. Hartmann, J. Keller, T. Neubert, “A Revised Taxonomy of Steganography Embedding Patterns”, in Proc. of the 16th International Conference on Availability, Reliability and Security, Vienna, Austria, pp. 67:1 – 67:12, Aug. 2021

Open Research Challenges (4)

- Watermarking is useful to give an **advantage** to defendants
- Part of the **ongoing research** tries to hide information to identify:
 - AI-generated contents
 - fake news and deep fakes
 - large-scale ingestion of copyrighted data
- For the specific case of LLMs:
 - very efficient in generating code
 - used to create multiple (unbounded) mutations of **malicious** routines
 - this can challenge signature-based detection mechanisms
 - watermarking LLMs-generated code could restore a balance!



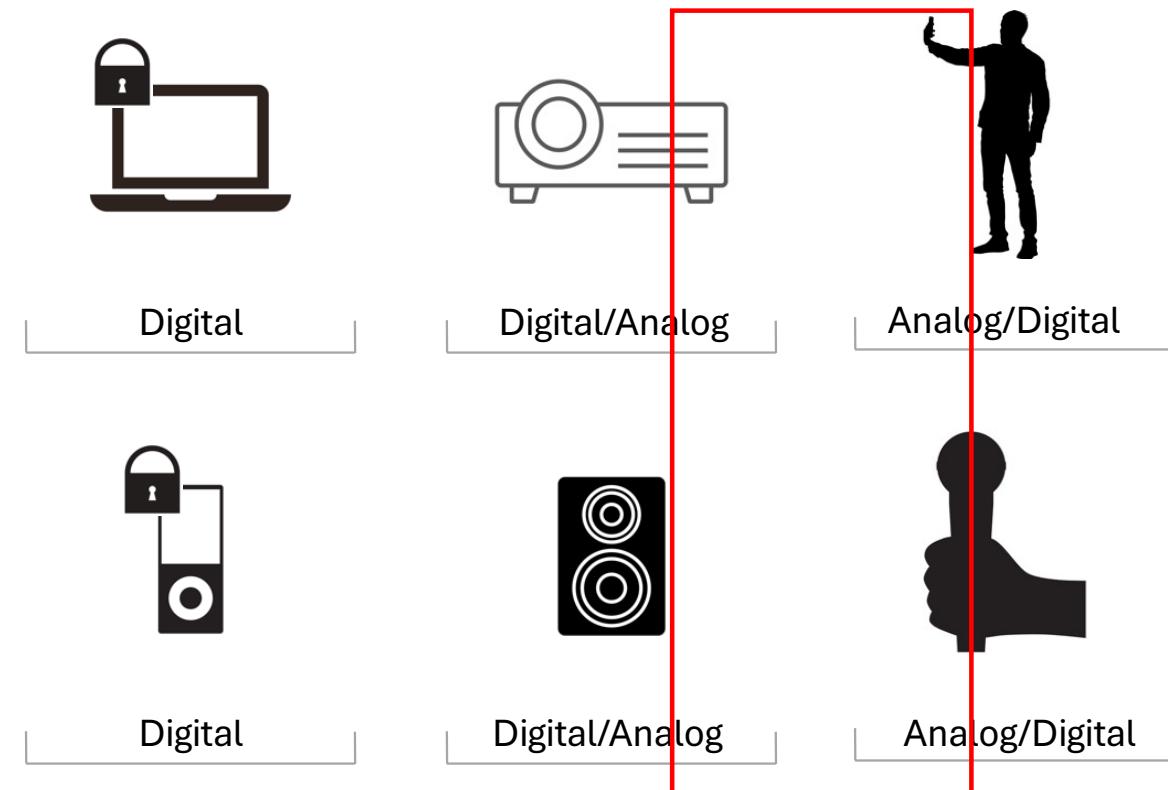
Watermark (e.g., opaque predicate)

Open Research Challenges (5)

- Recommendations and regulations(*):
 - mostly on a **country-by-country** basis
 - nowadays a vivid area of **discussion** and research
- Some examples:
 - **EU** Artificial Intelligence Act: recommendations on watermarking contents
 - **US** - Ensuring Safe, Secure and Trustworthy AI: guidelines for watermarking audio and visual contents
 - **China**: many AI services should display proper information to make users aware of the content creator (e.g., human vs machine)
 - **Microsoft**: announced watermarking capabilities for their tools, especially to fight diffusion of fake contents

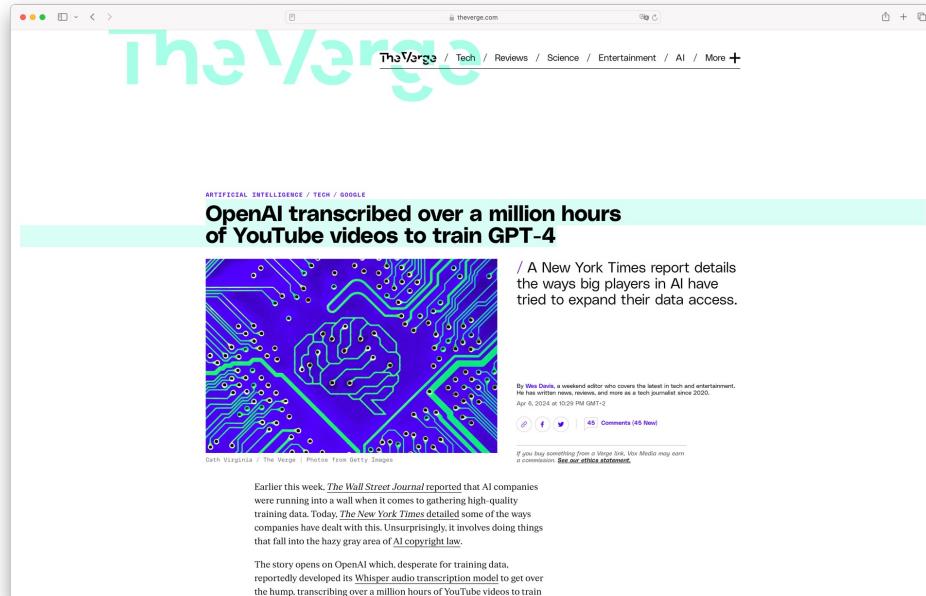
Open Research Challenges (6)

- Digital watermarking is susceptible to **analog holes**

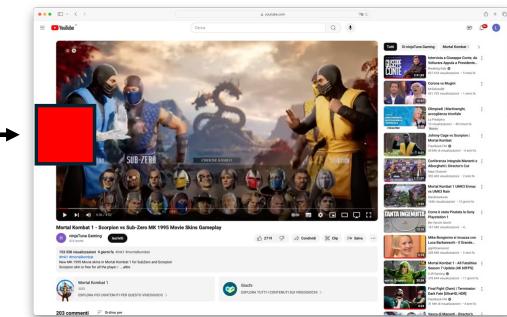


Open Research Challenges (6)

- Cross-domain watermark for facing new holes?



Watermark



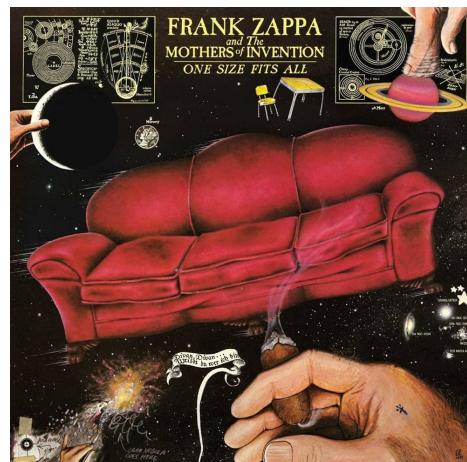
Lore [REDACTED] ipsum dolor sit amet, consectetur adipisicing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo cons[REDACTED]t. Duis aute irure dolor in reprehenderit in voluptate velit e[REDACTED]ccillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum. Sed ut perspiciatis unde omnis iste natus error sit voluptatem accusantium [REDACTED] premque laudantium, totam rem aperiam, eaque ipsa quae ab illi d[REDACTED]entore veritatis et quasi architecto beatae vitae

OpenAI transcribed over a million hours of YouTube videos to train GPT-4: <https://www.theverge.com/2024/4/6/24122915/openai-youtube-transcripts-gpt-4-training-data-google>

Takeaways

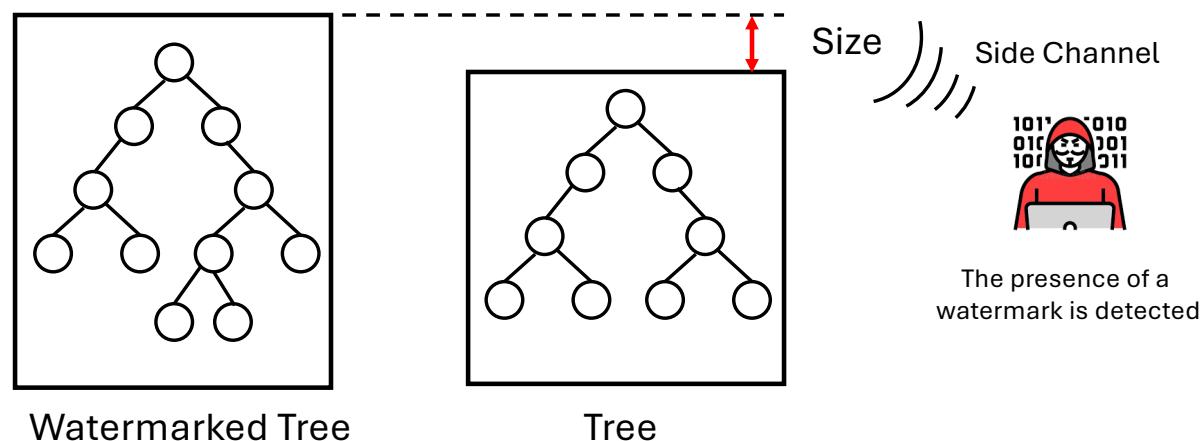
Takeaways (1)

- Watermarking is tightly-coupled with:
 - **the carrier:** there is no a unique hiding strategy
 - **the application:** how data will be used/retained plays a major role
 - **the processing pipeline:** it may disrupt the information
 - **the attack model:** each threat requires a specific solution



Takeaways (2)

- Double-edged cases everywhere:
 - the need of being sometimes **robust**, sometimes **fragile**
 - offering solutions to **unwanted tracking** – profiling
 - ability of an attacker of “piggyback” a method to **conceal information**
 - **side-channels** and **data-leakages**



Takeaways (3)

- Watermarking is a very **heterogeneous** area borrowing ideas from:
 - steganography and steganalysis
 - information hiding and software/cyber security
 - overlapping and “incompatible” fields

A cloud of colored words representing concepts related to watermarking, including:

- poisoning
- simplification
- copyright attacks
- frequency
- black-box
- information-hiding
- robust lsb
- positional
- secrecy
- integrity
- visible
- capacity
- noise
- decoding
- efficiency
- robustness
- steganography
- encryption
- crop
- size
- tracking
- fidelity
- generality
- embedder
- decoder
- timing
- spatial
- reliability
- invisible
- content
- carrier
- white-box
- watermarking
- fragile
- backdoor
- side-channel
- encoding
- secret-key
- tampering
- compression
- profiling
- substitution
- smoothing

Conclusions

Conclusions (1)

- Watermarking can be used for:
 - **defense**, e.g., to track or protect software assets and AI models
 - **attack**, e.g., to hide malicious data or create covert communication paths
- No **one-size-fits-all** mechanisms for:
 - the embedding process
 - attacks
 - applications
 - design of countermeasures against removal (not covered in this class)
- Data can be hidden (almost) in any media or digital object

Conclusions (2)

- **Information hiding** is a vivid research area expected to **grow**:
 - driven by **AI** (and satellites areas)
 - to tame the **impact** of data and technology over **people**
- Challenges in watermarking are not limited to:
 - cybersecurity
 - regulations
 - performances
- **We only scratched the surface!**
- Thank you for listening
- Questions?